How NOT to Create an Um Detector Using Create ML

Yono Mittlefehldt — @yonomitt — 2022

Disclaimer

Disclaimer

Live Demo

Live Demo

Listening to live audio

Running a Core ML model

Analyzing for “um”s

Making predictions every 0.5s

Live Demo

Listening to live audio

Running a Core ML model

Analyzing for “um”s

Making predictions every 0.5s

…

Counting my shame

Who is Yono?

Co-creator of Gus on the Go

Who is Yono?

Co-creator of Gus on the Go

Freelance software developer
specializing in Computer Vision and iOS

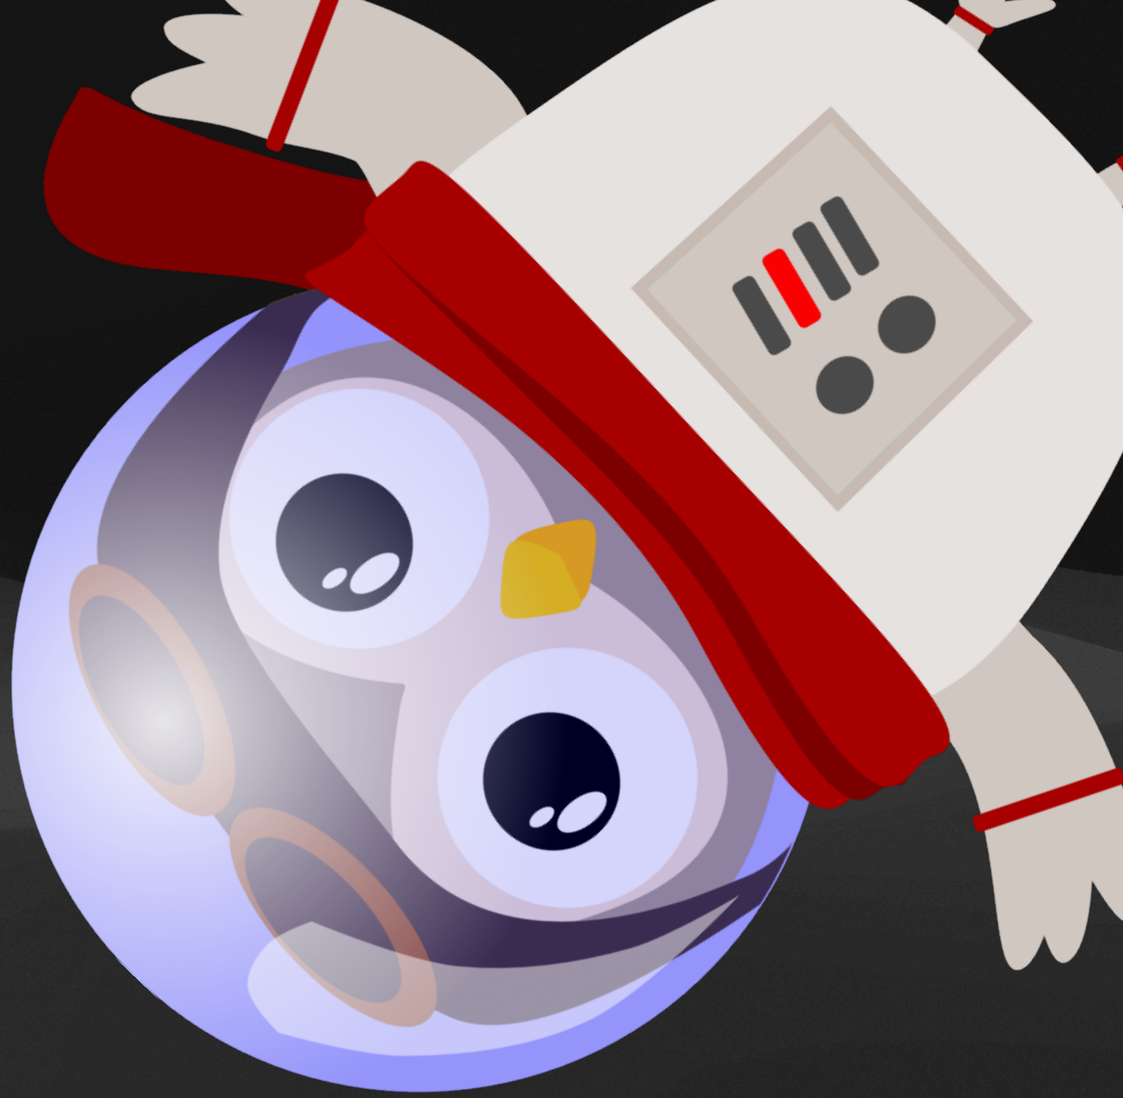

Who is Yono?

Co-creator of Gus on the Go

Freelance software developer specializing in Computer Vision and iOS

Tutorial writer for raywenderlich.com

Who is Yono?

Co-creator of Gus on the Go

Freelance software developer specializing in Computer Vision and iOS

Tutorial writer for raywenderlich.com

As a child, I once got to operate the volcano on a float in a NASA parade

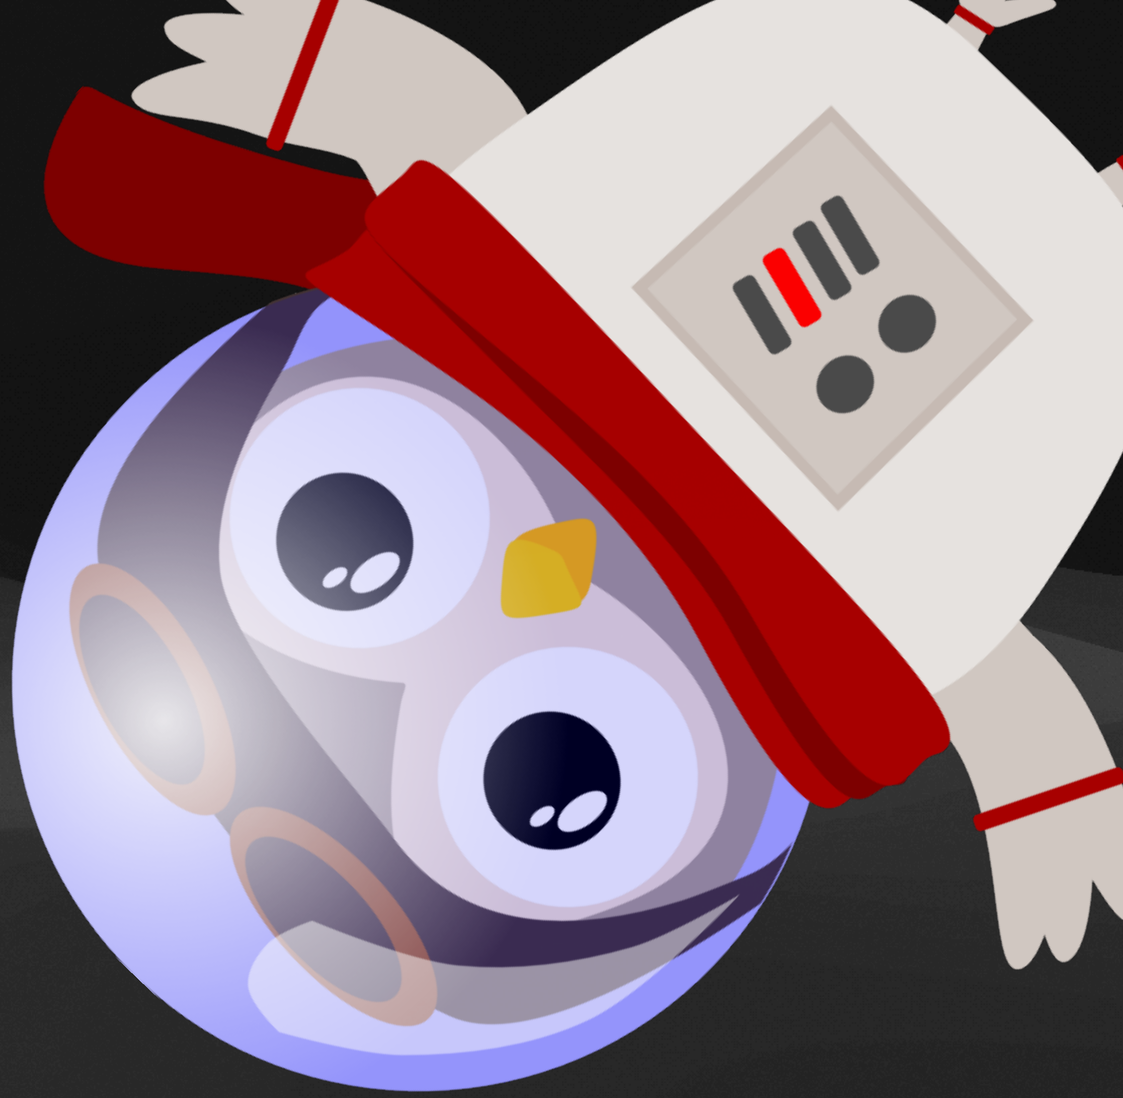

TREK the next generation

Story Time

Why would you want to do this?

Who needs this?

Podcasters

Audio editors

Public speaking trainers

Who needs this?

Podcasters

Audio editors

Public speaking trainers

…

YouTubers?

This talk isn’t about code

It’s really about a thought process

How do we detect ums?

Attempt #1 - Speech Recognition

Speech Recognition

Apple’s Speech Framework

It’s built in

It’s easy to use

Falls under the umbrella of Apple’s privacy policy

```swift
let audioEngine = AVAudioEngine()
```

```swift
// Configure the microphone input
let audioEngine = AVAudioEngine()
let inputNode = audioEngine.inputNode
```

```swift
// Configure the microphone input
let audioEngine = AVAudioEngine()
let inputNode = audioEngine.inputNode

// Get the format of the data
let format = inputNode.outputFormat(forBus: 0)
```

AVAudioInputNode

AVAudioInputNode

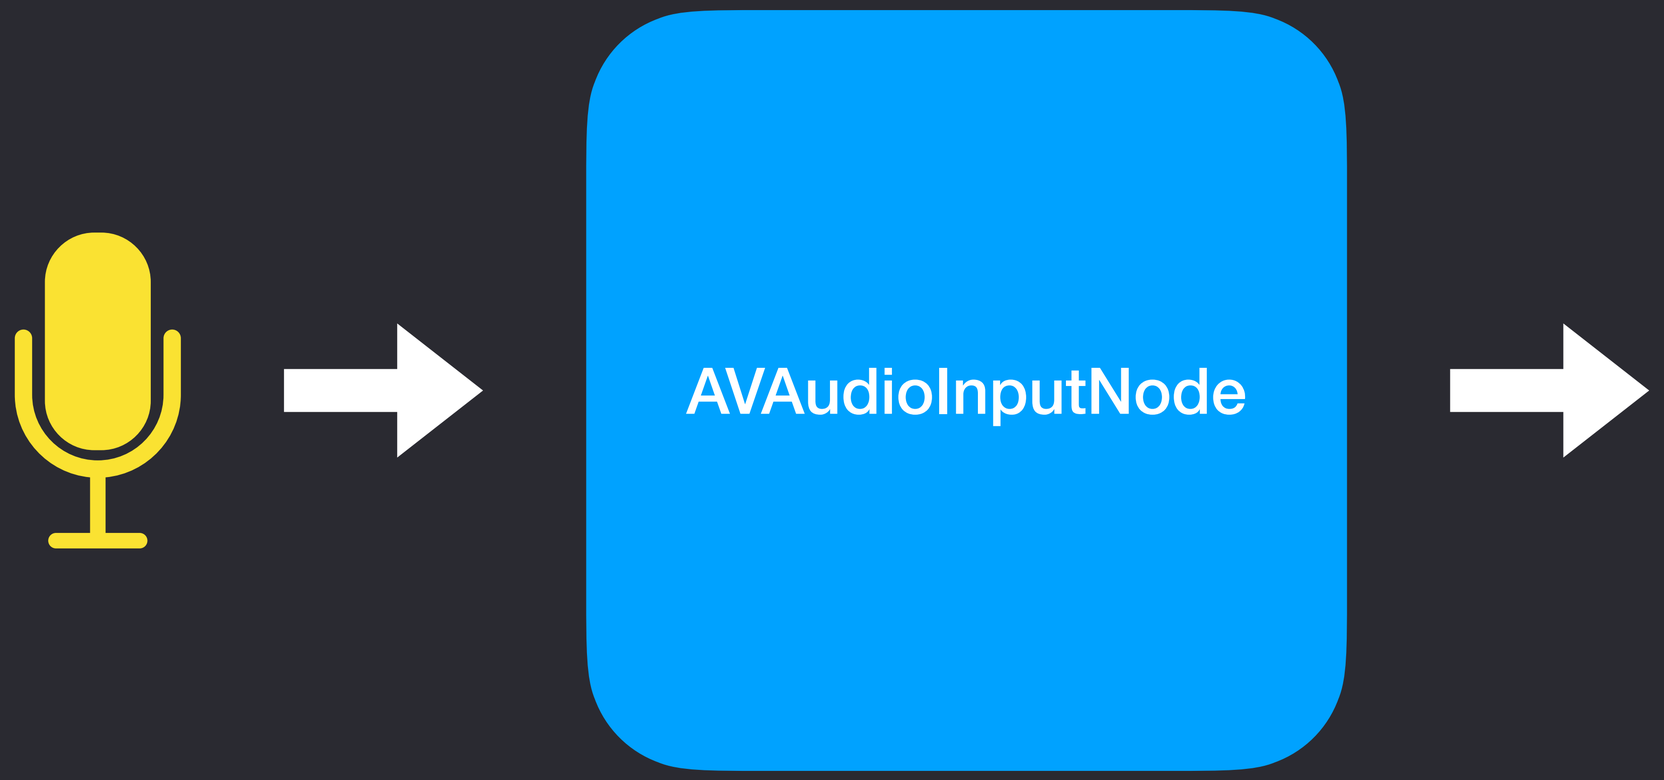
AVAudioInputNode

iPhone XS Max

```swift
// Configure the microphone input
let audioEngine = AVAudioEngine()
let inputNode = audioEngine.inputNode

// Get the format of the data
let format = inputNode.outputFormat(forBus: 0)
```

```swift
// Configure the microphone input
let audioEngine = AVAudioEngine()
let inputNode = audioEngine.inputNode

// Get the format of the data
let format = inputNode.outputFormat(forBus: 0)

// Install an audio tap to observe the output of the input node
inputNode.installTap(onBus: 0, bufferSize: 1024, format: format) { (buffer, _) in
    request.append(buffer)
}
```

```swift
// Create the Speech request
let request = SFSpeechAudioBufferRecognitionRequest()

// Configure the microphone input
let audioEngine = AVAudioEngine()
let inputNode = audioEngine.inputNode

// Get the format of the data
let format = inputNode.outputFormat(forBus: 0)

// Install an audio tap to observe the output of the input node
inputNode.installTap(onBus: 0, bufferSize: 1024, format: format) { (buffer, _) in
    request.append(buffer)
}
```

But it doesn’t work

How not to create
an arm detector
using create ML

How not to create
an M detector

Hacky Hack

Um Contact

Allows dictation to detect “um”s

Sorta…

Only detected if “um” could be replaced by a name

“Uh” didn’t work

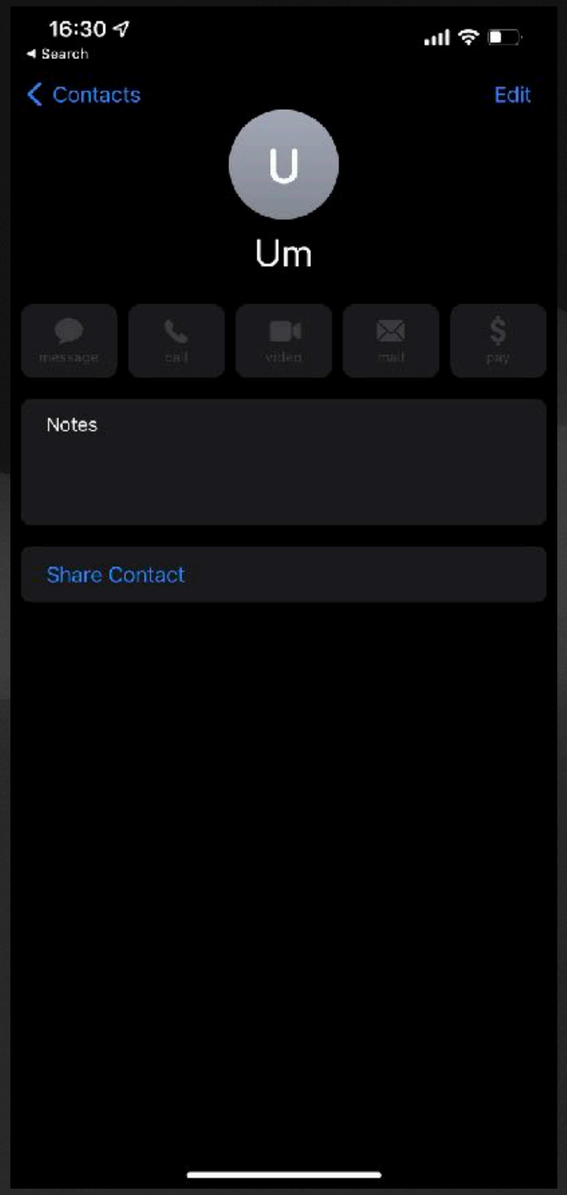

How do we detect ums?

Attempt #2 - Machine Learning

“Machine Learning is the science of getting computers to learn without being explicitly programmed.”

Sebastian Thrun

X
Y
A
B

Chat Bot User Interface
Temp Name
Temp Data

Machine Learning

What do we need?

Lots of math

Dataset

Train a model

Testing

Machine Learning

What do we need?

~~Lots of math~~

Dataset

Train a model

Testing

Dataset

Possibly the most important part

Raw audio files

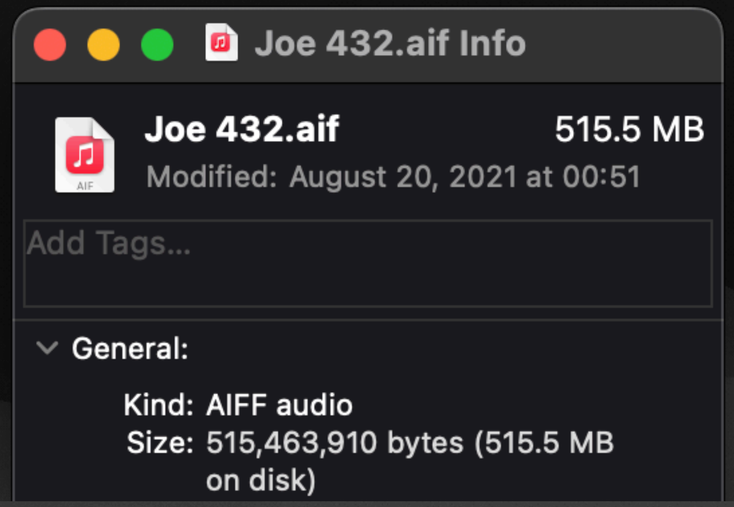

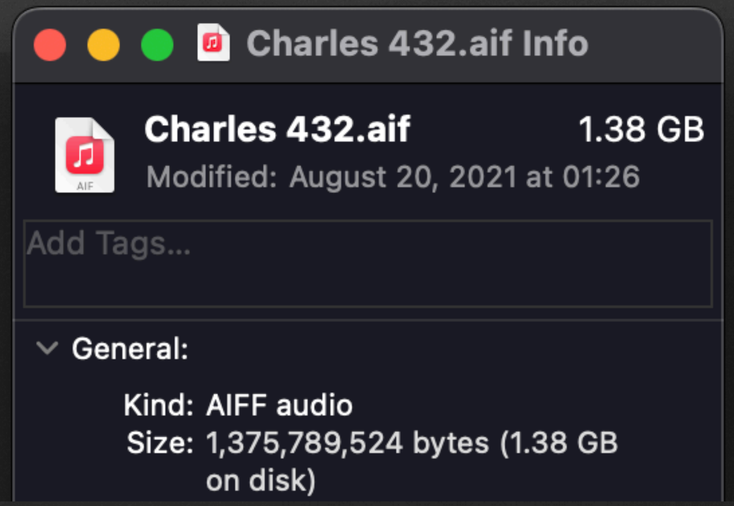

Dataset

Possibly the most important part

Raw audio files

Cut them into 1 second segments

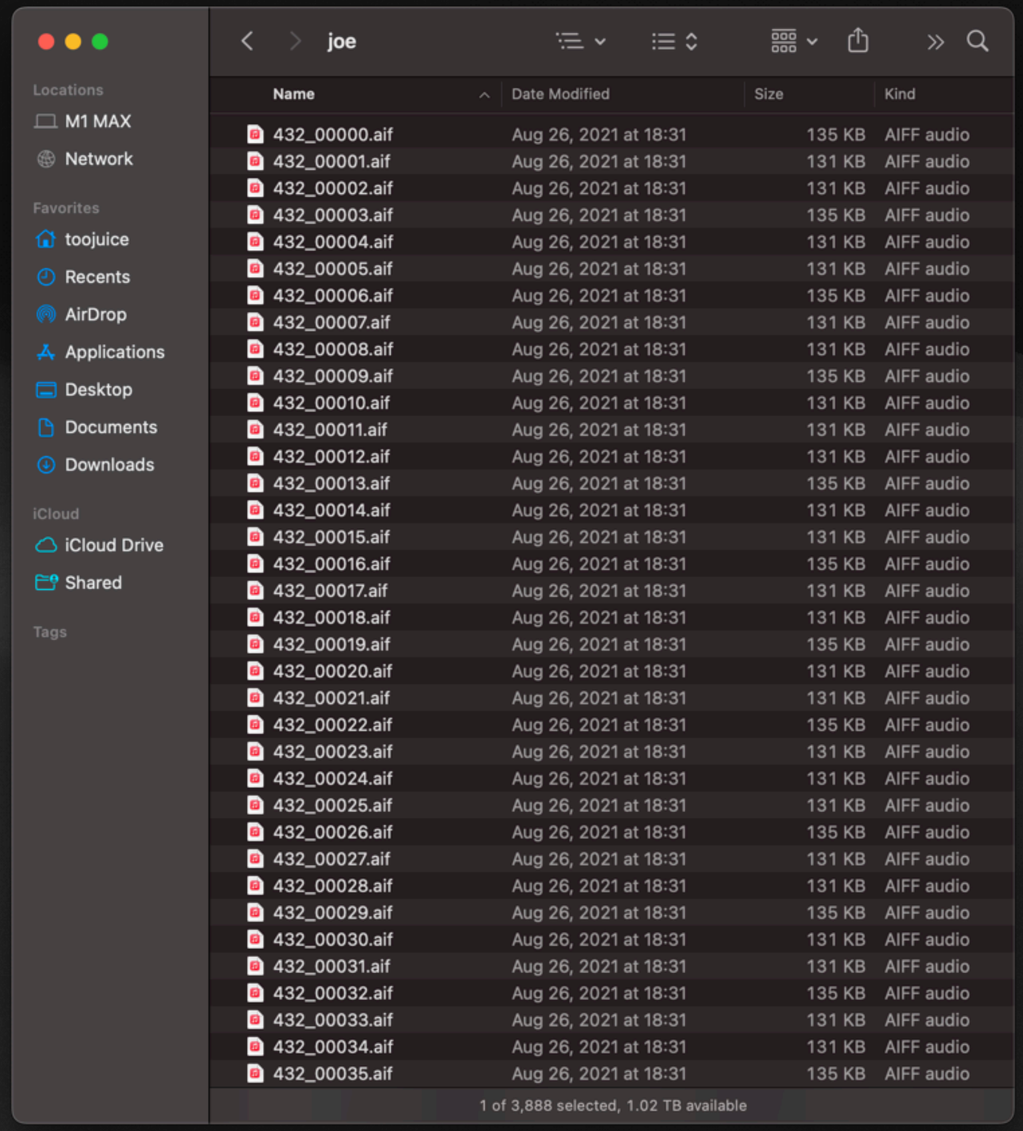

```
> ffmpeg -i Joe\ 432.aif -f segment -segment_time 1 -c copy joe/432_%05d.aif
```

```
ffmpeg           – the command
-i Joe \432.aif  – input file
-f segment       – segment the file
-segment_time 1  – 1 second segments
-c copy          – copy the stream, don’t transcode
Joe/432_%05d.aif – output files should be numbered with 5 digits
```

Dataset

Possibly the most important part

Raw audio files

Cut them into 1 second segments

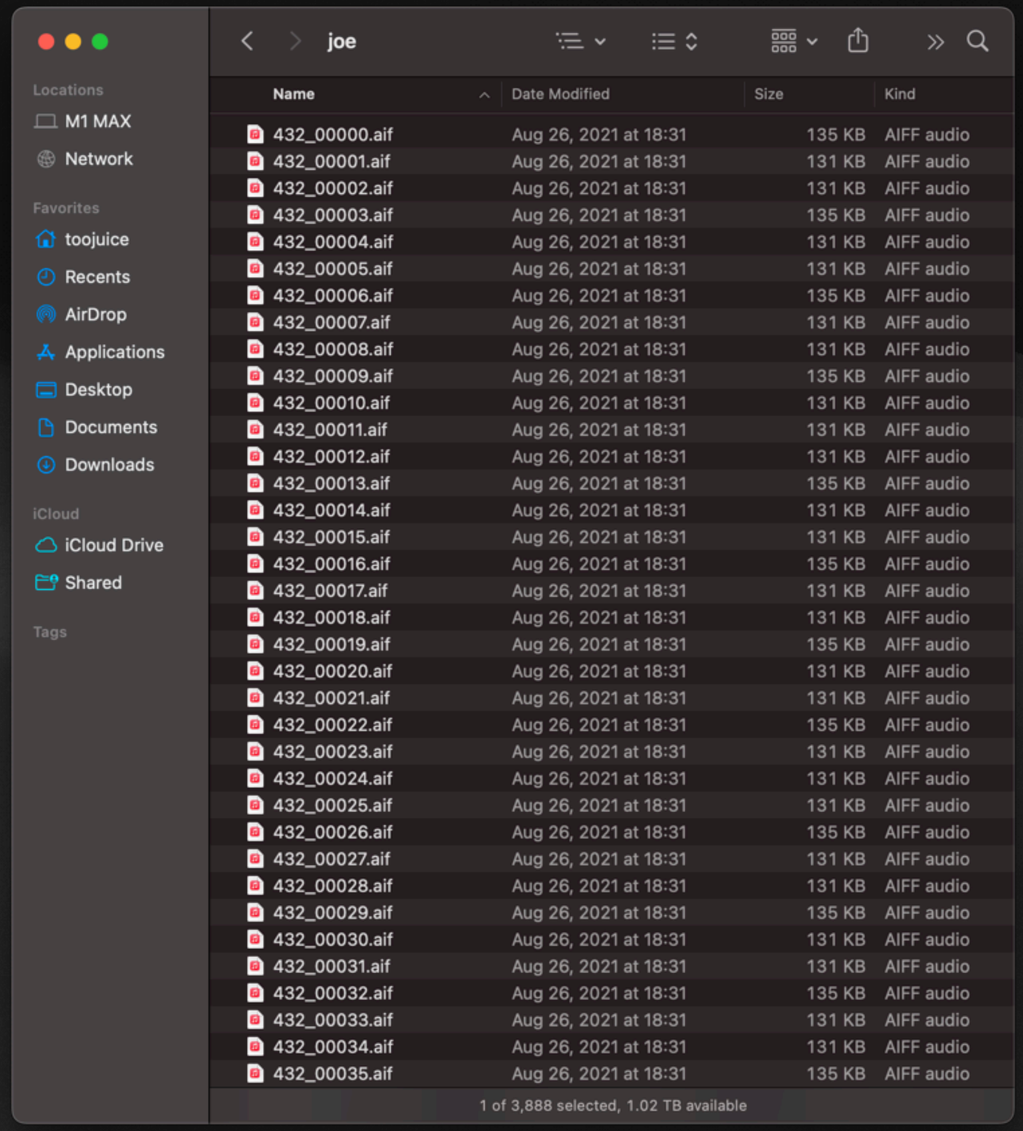

Dataset

Possibly the most important part

Raw audio files

Cut them into 1 second segments

Sort them into categories

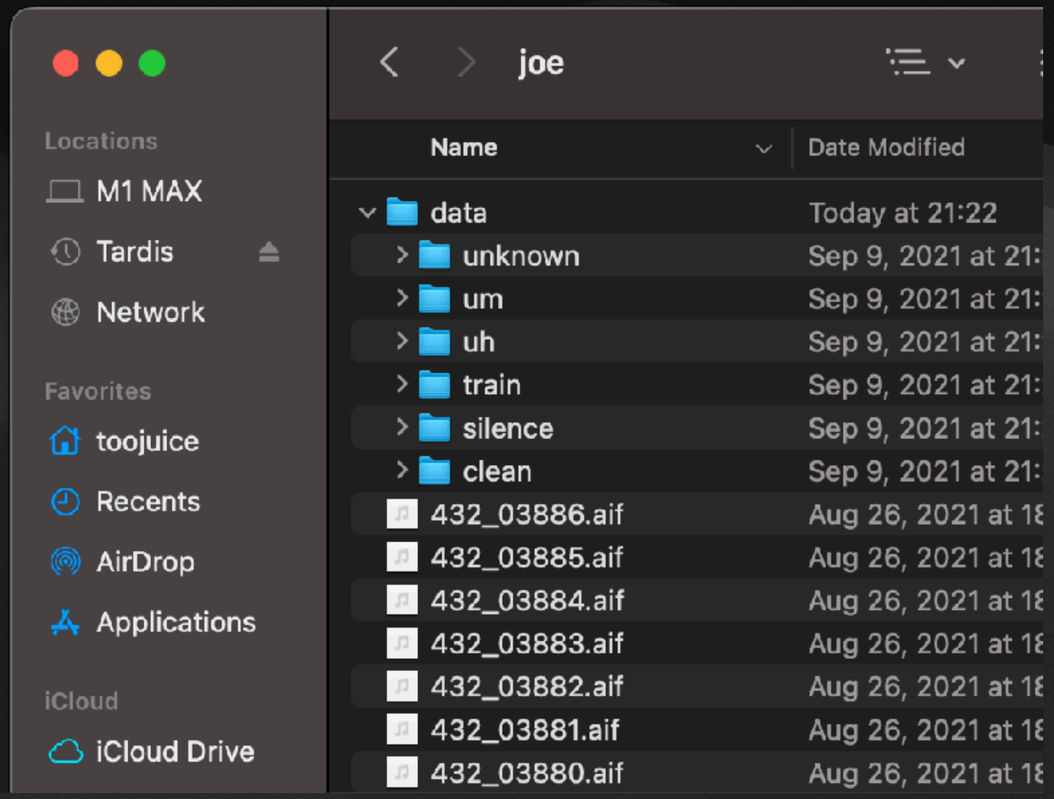

Clean

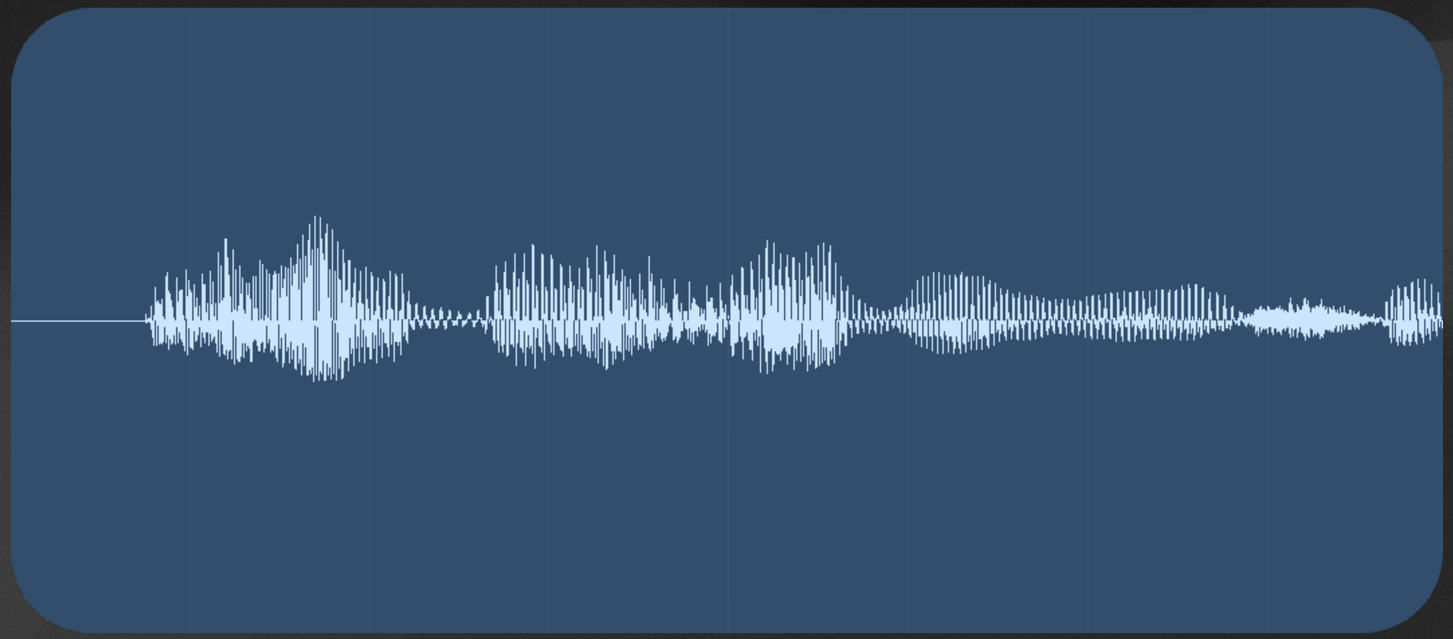

Um

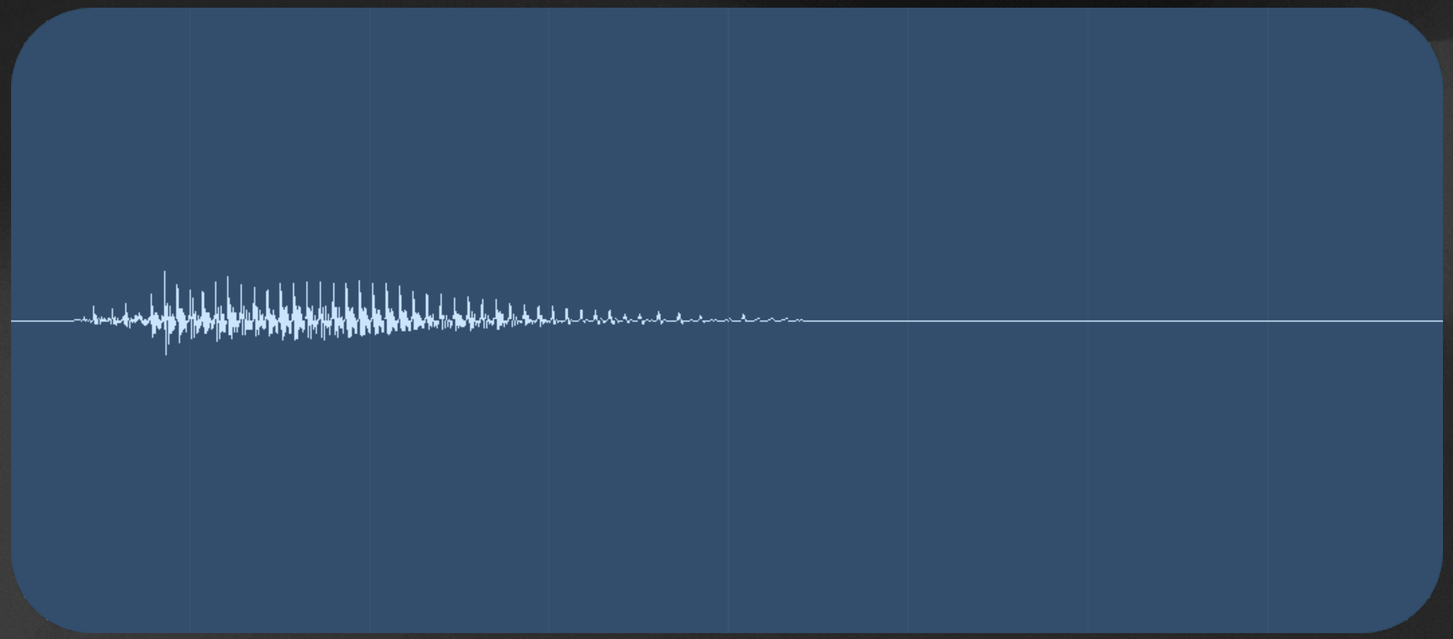

Dataset

Possibly the most important part

Raw audio files

Cut them into 1 second segments

Sort them into categories

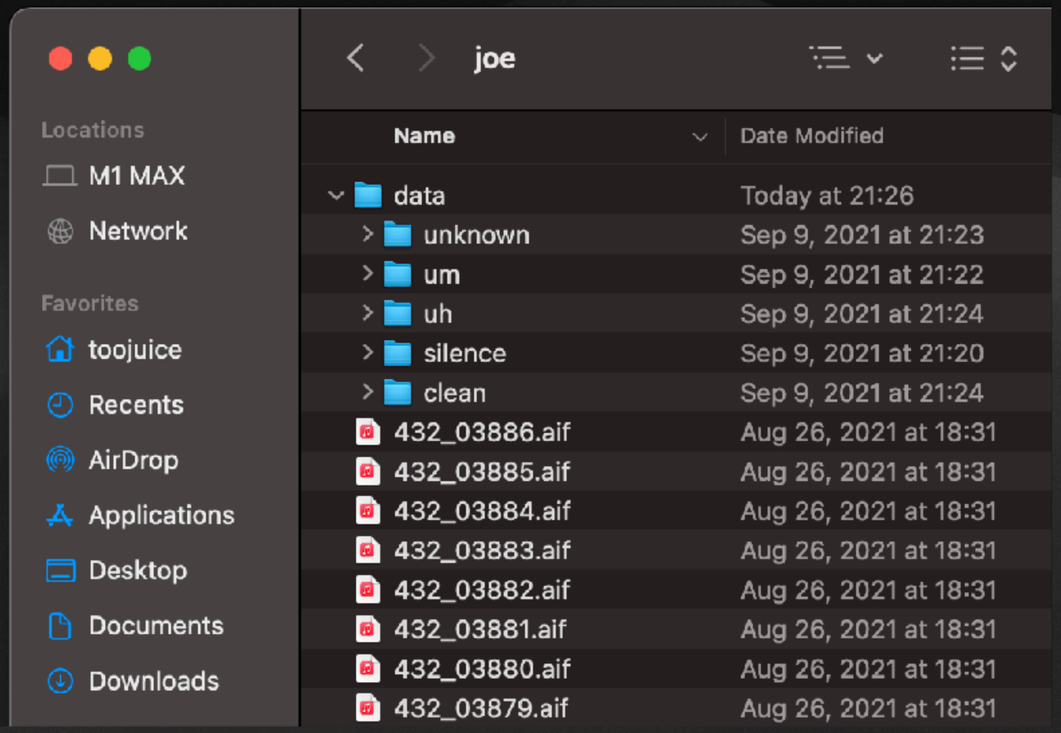

Silence

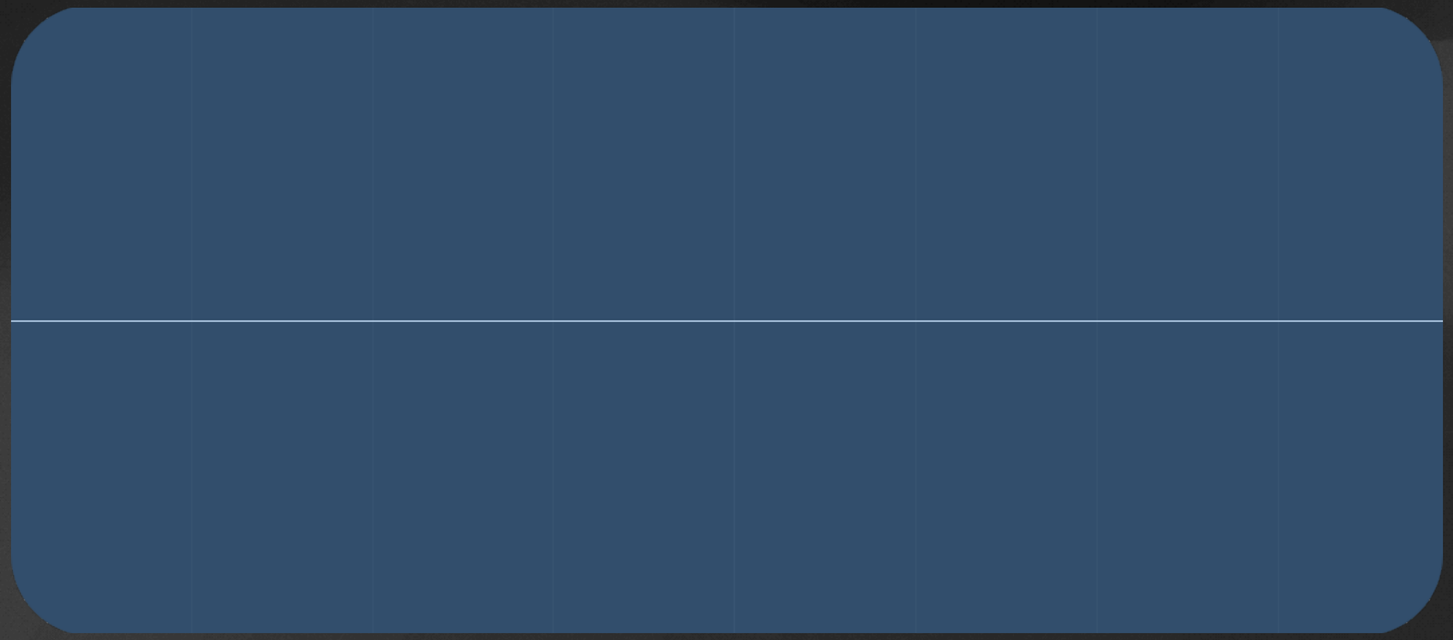

Uh

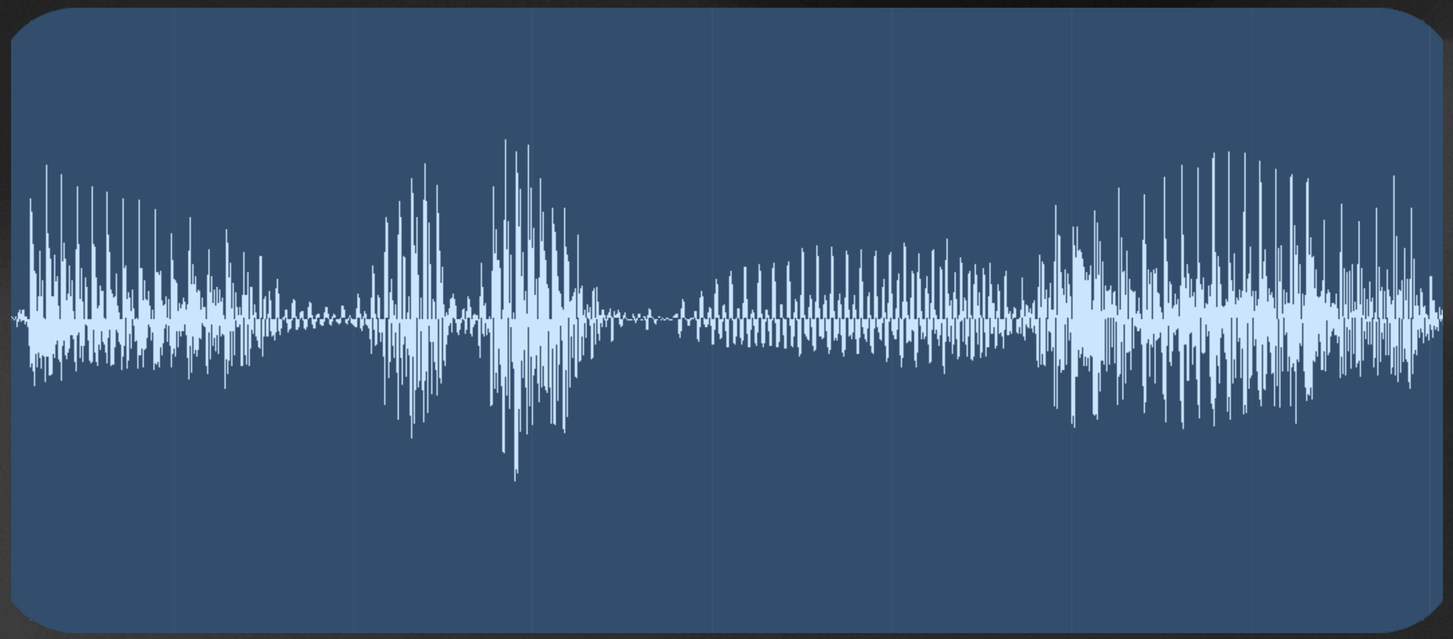

Unknown

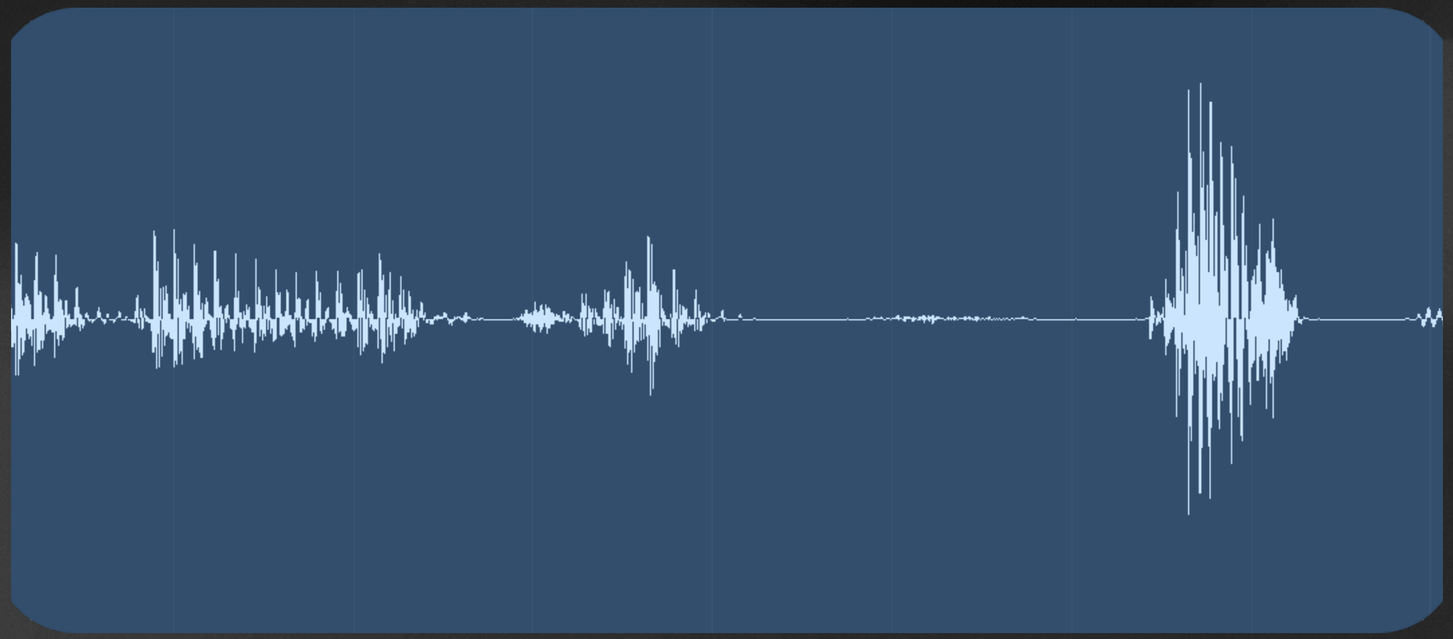

“This seems tedious”

You, probably

Touch Bar

Dataset

Possibly the most important part

Raw audio files

Cut them into 1 second segments

Sort them into categories

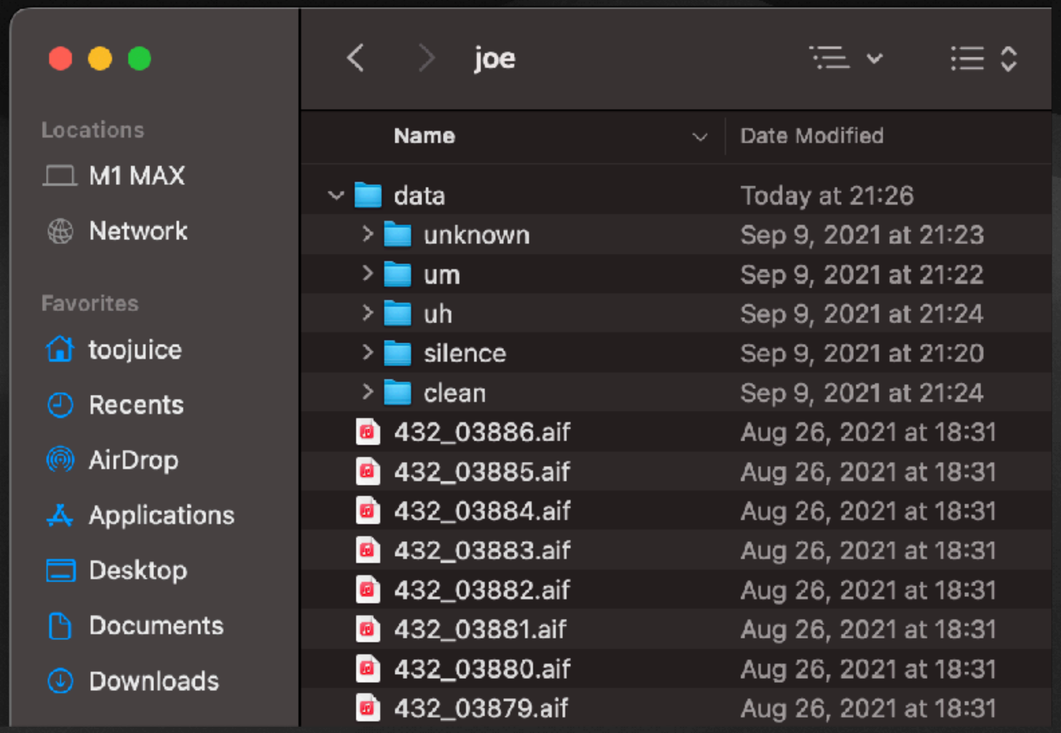

Dataset

Possibly the most important part

Raw audio files

Cut them into 1 second segments

Sort them into categories

Leave some for testing

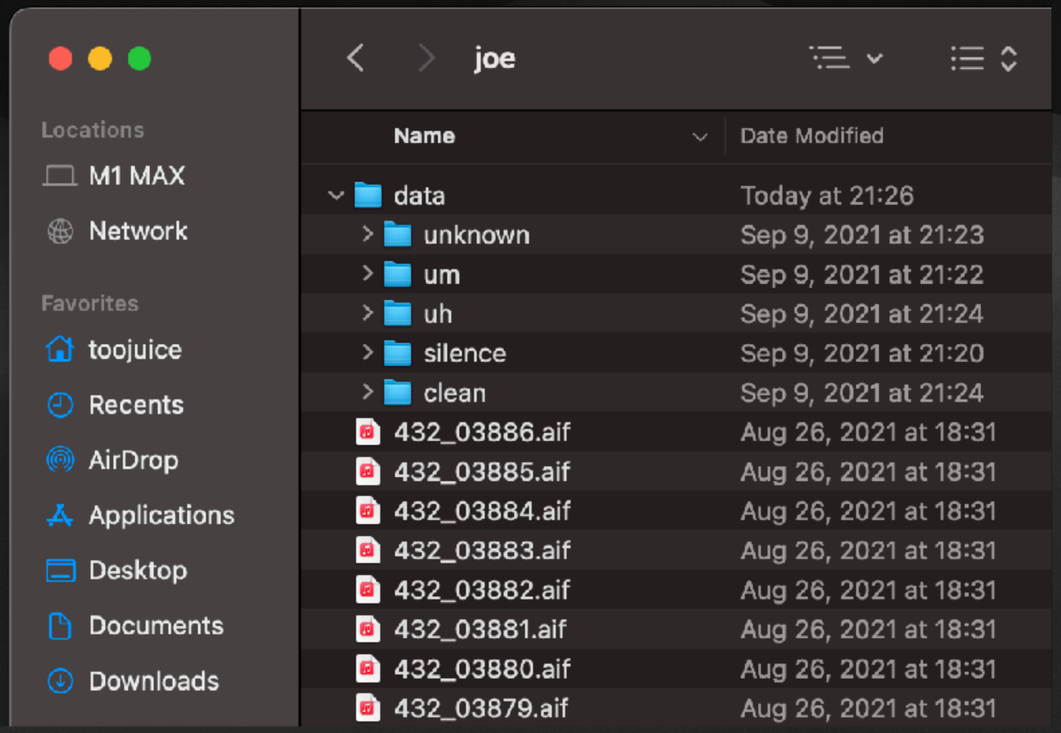

Machine Learning

What do we need?

~~Lots of math~~

Dataset

Train a model

Testing

Create ML

Create a sound classification project in Create ML

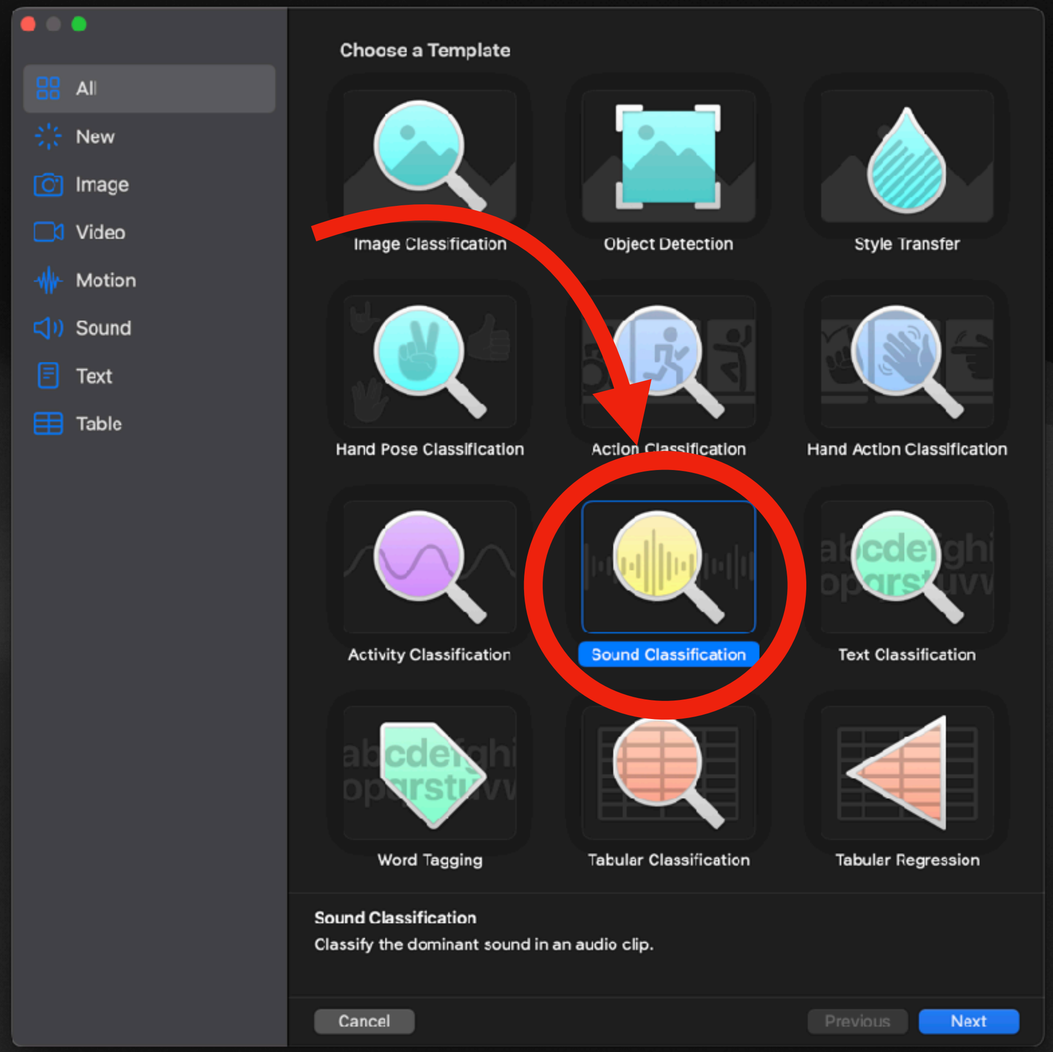

Create ML

Create a sound classification project in Create ML

Set project info

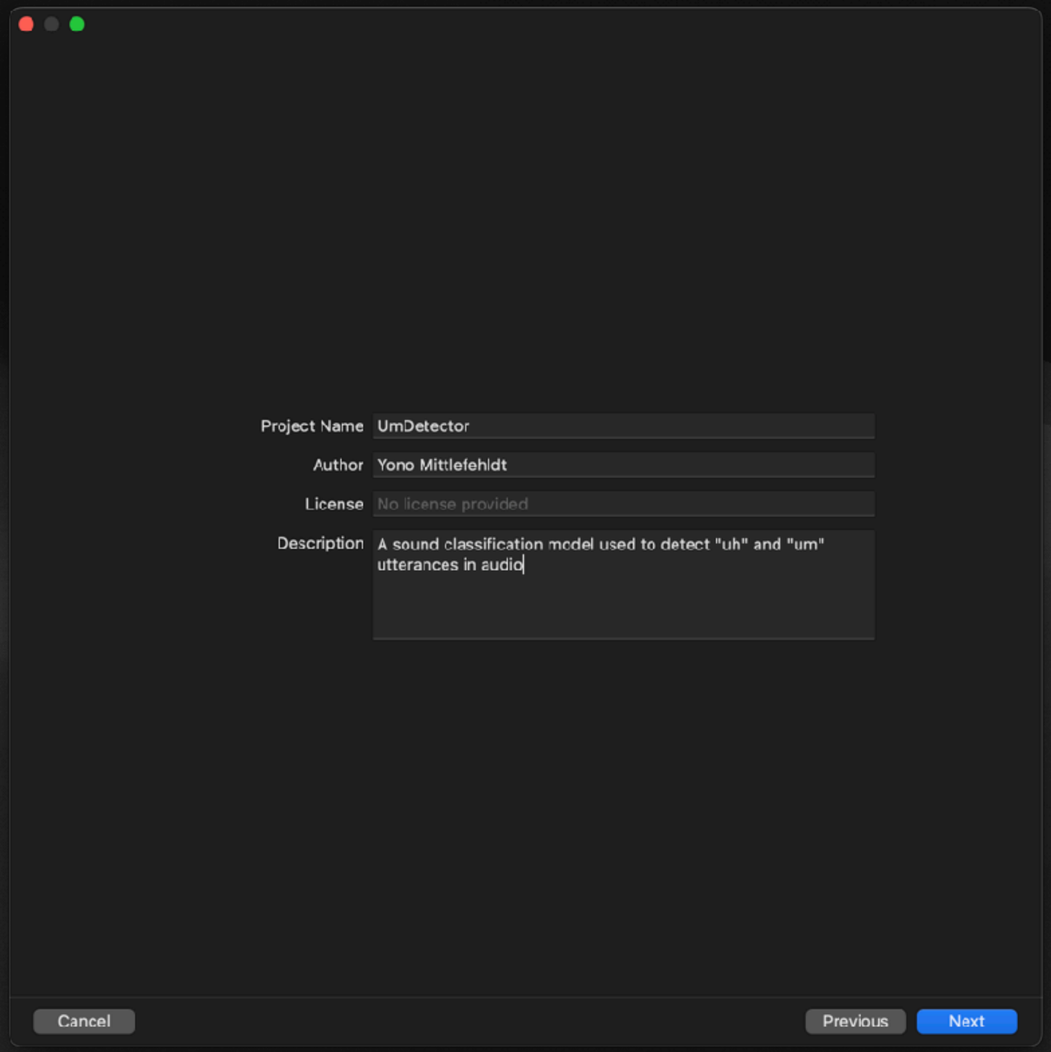

Create ML

Create a sound classification project in Create ML

Set project info

Add the dataset

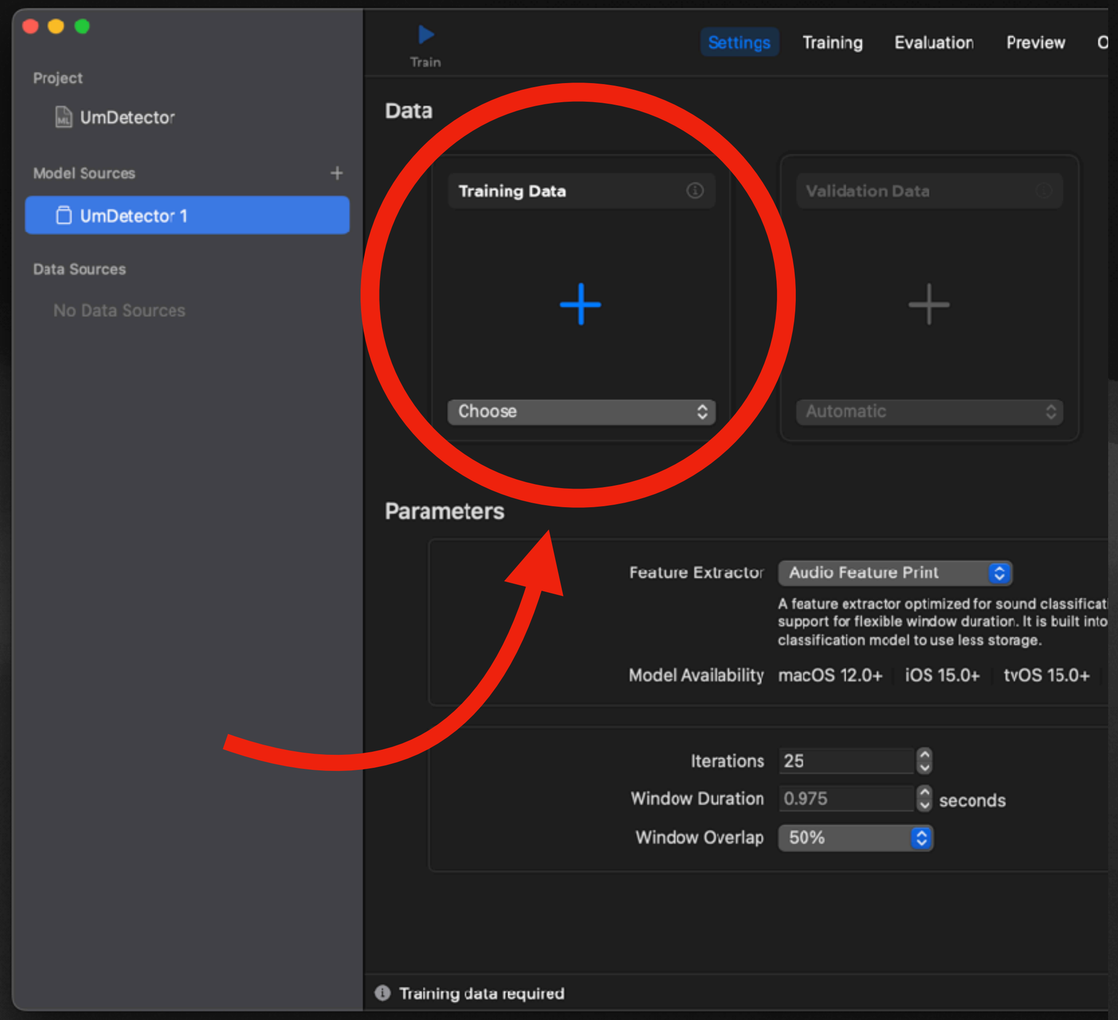

Dataset

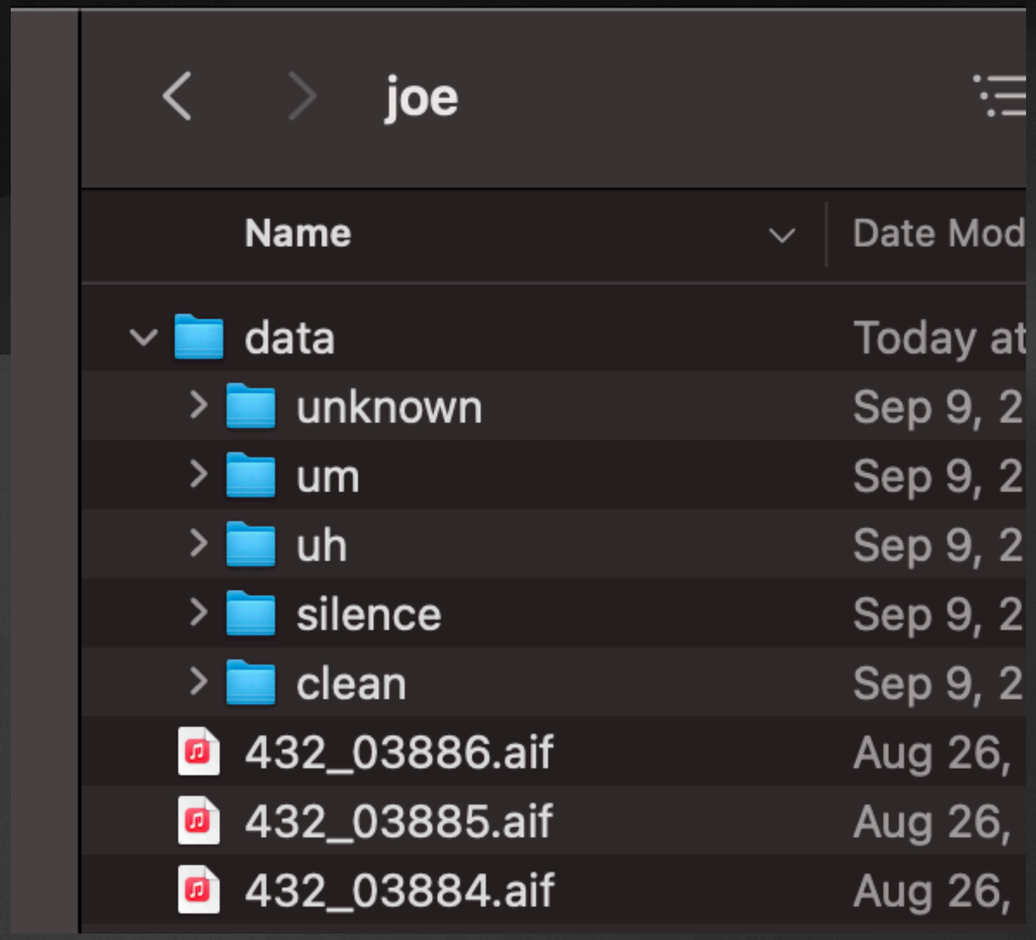
joe
Name
Date Mod
data
Today at
unknown
Sep 9, 2
um
Sep 9, 2
uh
Sep 9, 2
silence
Sep 9, 2
clean
Sep 9, 2
432_03886.aif
Aug 26,
432_03885.aif
Aug 26,
432_03884.aif
Aug 26,

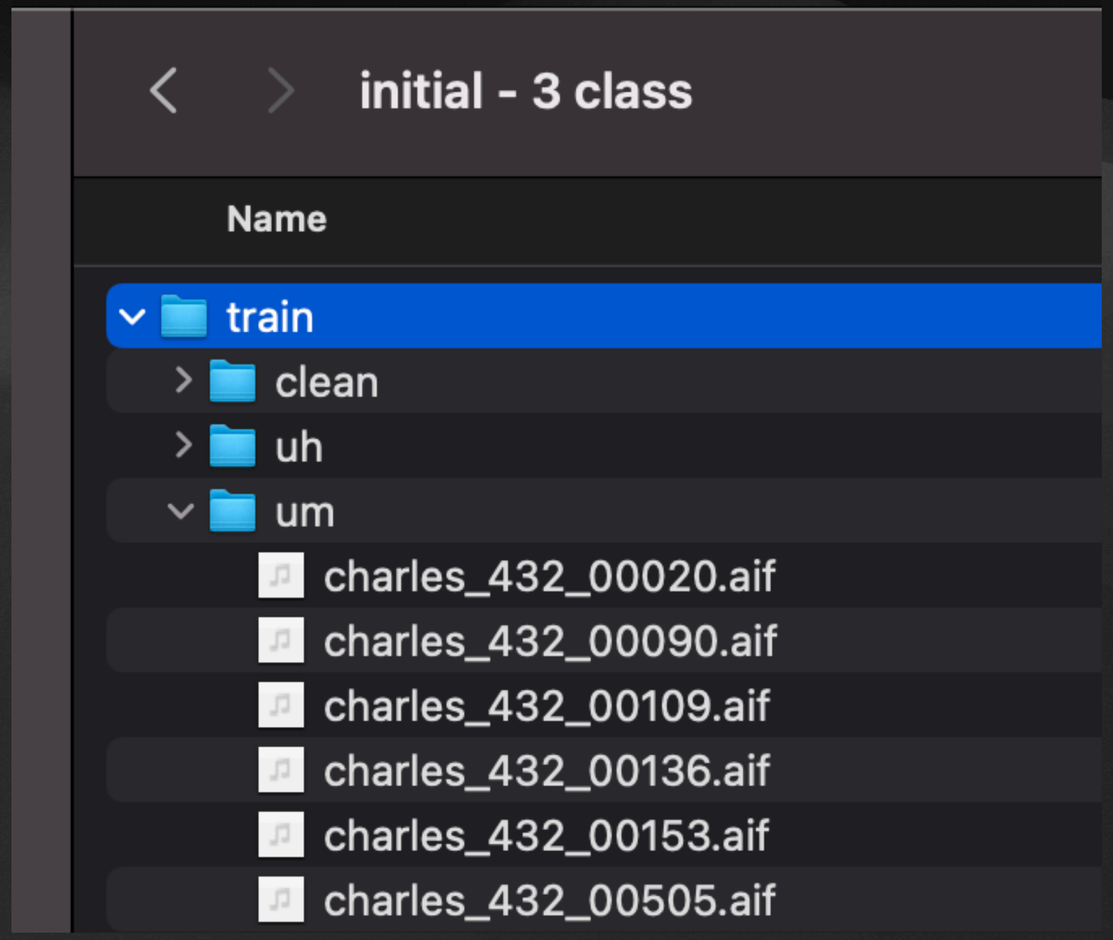
initial - 3 class
Name
train
clean
uh
um
charles_432_00020.aif
charles_432_00090.aif
charles_432_00109.aif
charles_432_00136.aif
charles_432_00153.aif
charles_432_00505.aif

Two vs Three Classes

Um vs Uh

Initially only two class

- Clean
- Um

Two vs Three Classes

Um vs Uh

Initially only two class

- Clean
- Um

Dataset creation unveiled 3rd class

- Uh

Dataset

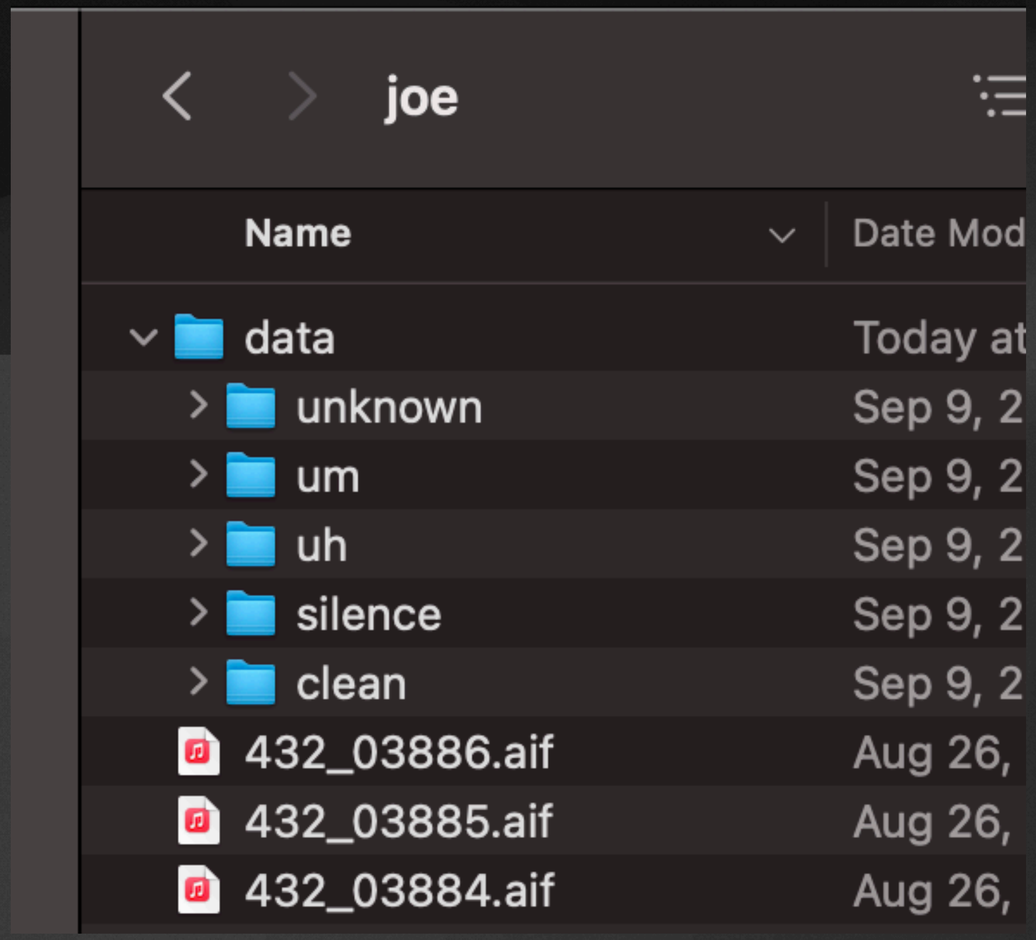
joe
Name
Date Mod
data
Today at
unknown
Sep 9, 2
um
Sep 9, 2
uh
Sep 9, 2
silence
Sep 9, 2
clean
Sep 9, 2
432_03886.aif
Aug 26,
432_03885.aif
Aug 26,
432_03884.aif
Aug 26,

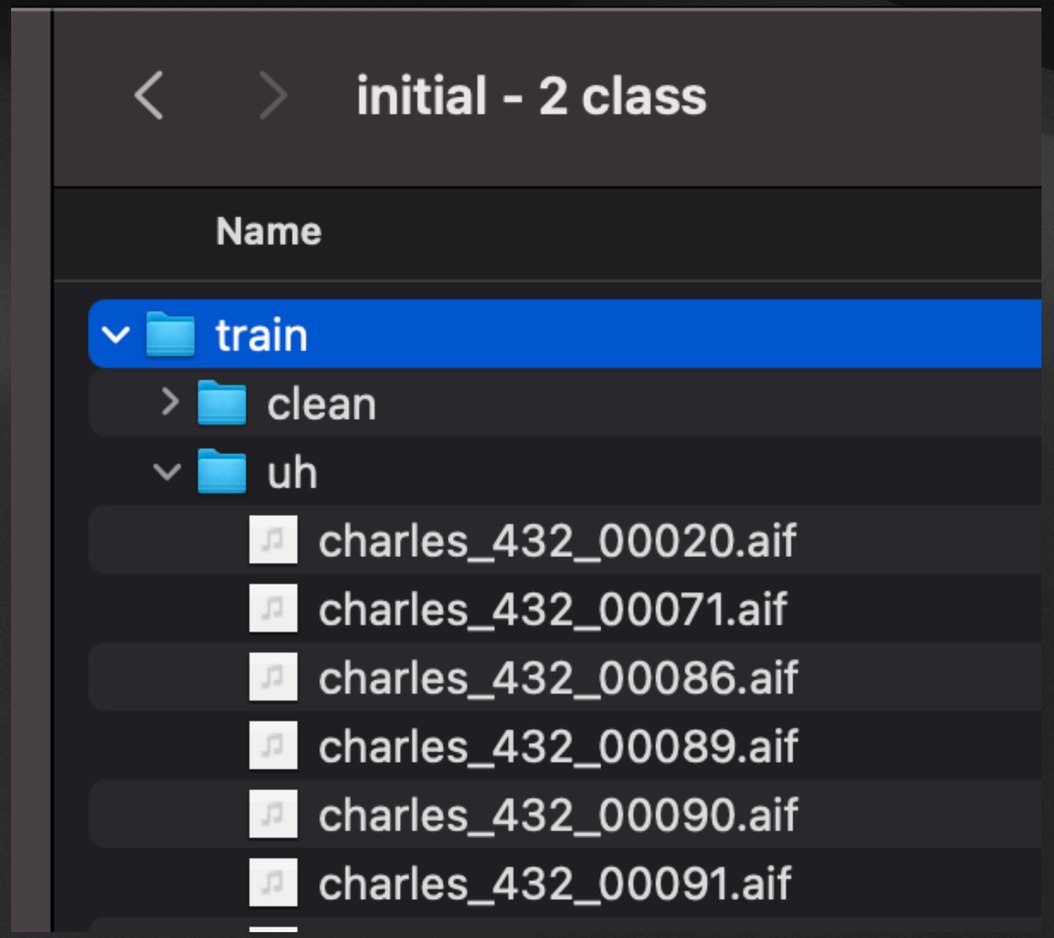
initial - 2 class
Name
train
clean
uh
charles_432_00020.aif
charles_432_00071.aif
charles_432_00086.aif
charles_432_00089.aif
charles_432_00090.aif
charles_432_00091.aif

Create ML

Create a sound classification project in Create ML

Set project info

Add the dataset

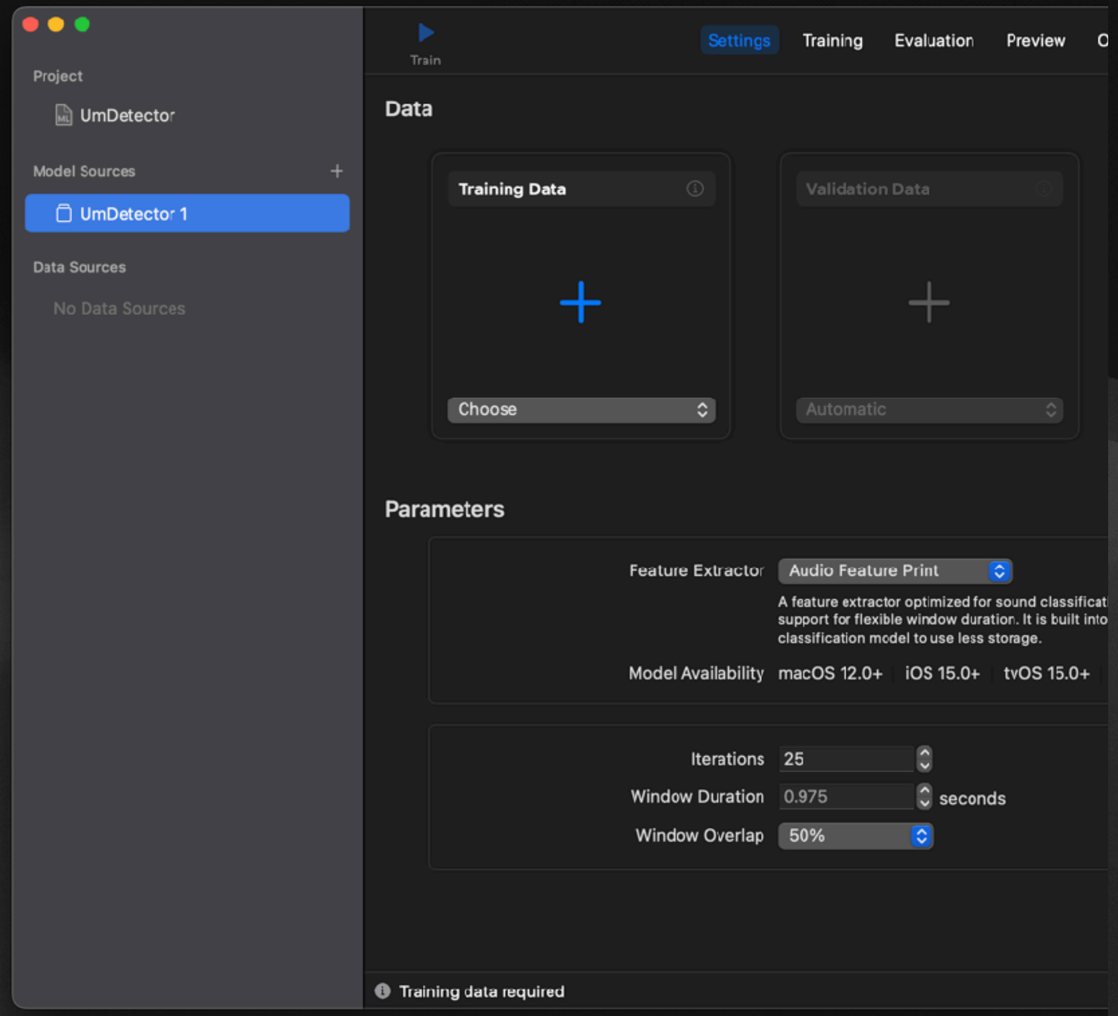

Create ML

Create a sound classification project in Create ML

Set project info

Add the dataset

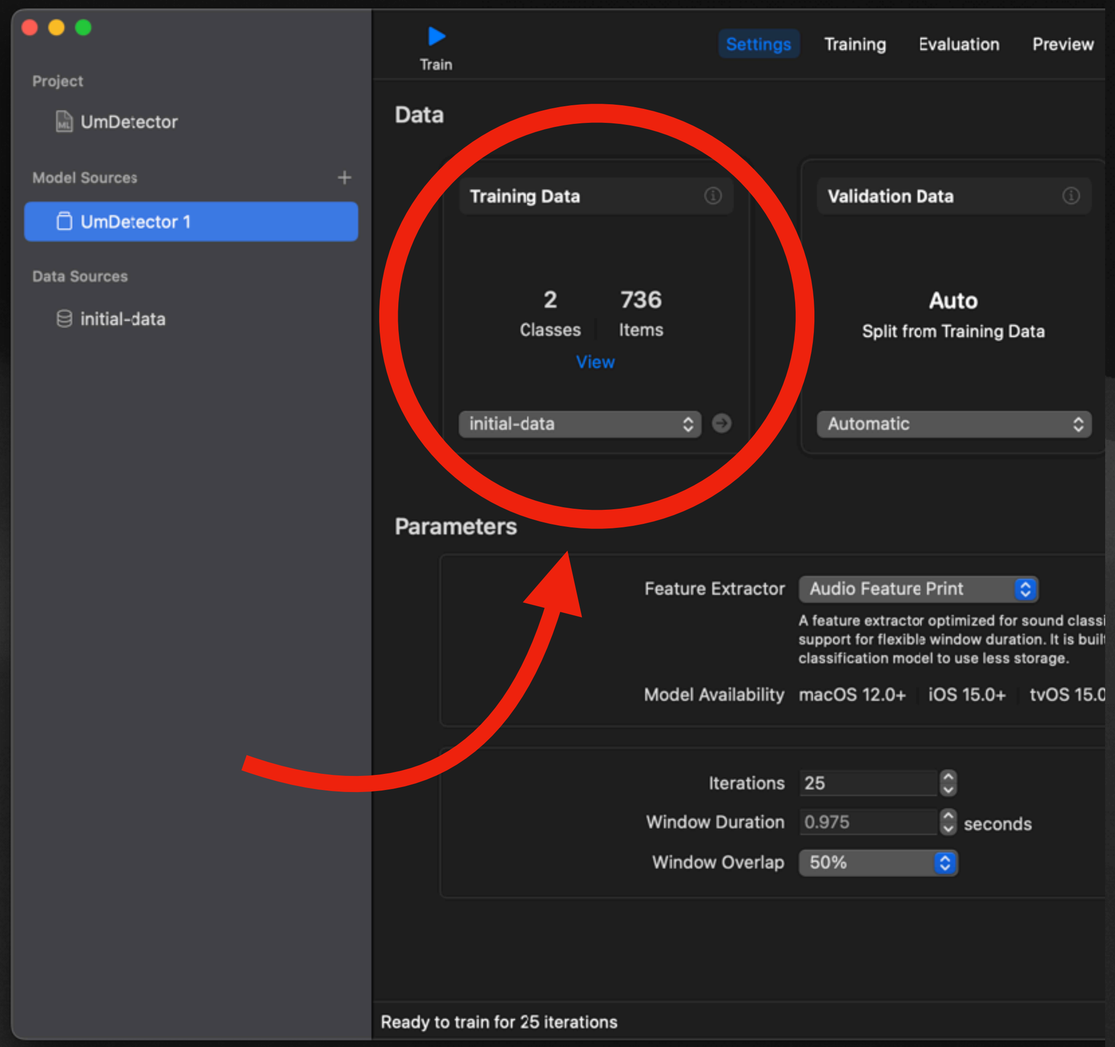

Create ML

Create a sound classification project in Create ML

Set project info

Add the dataset

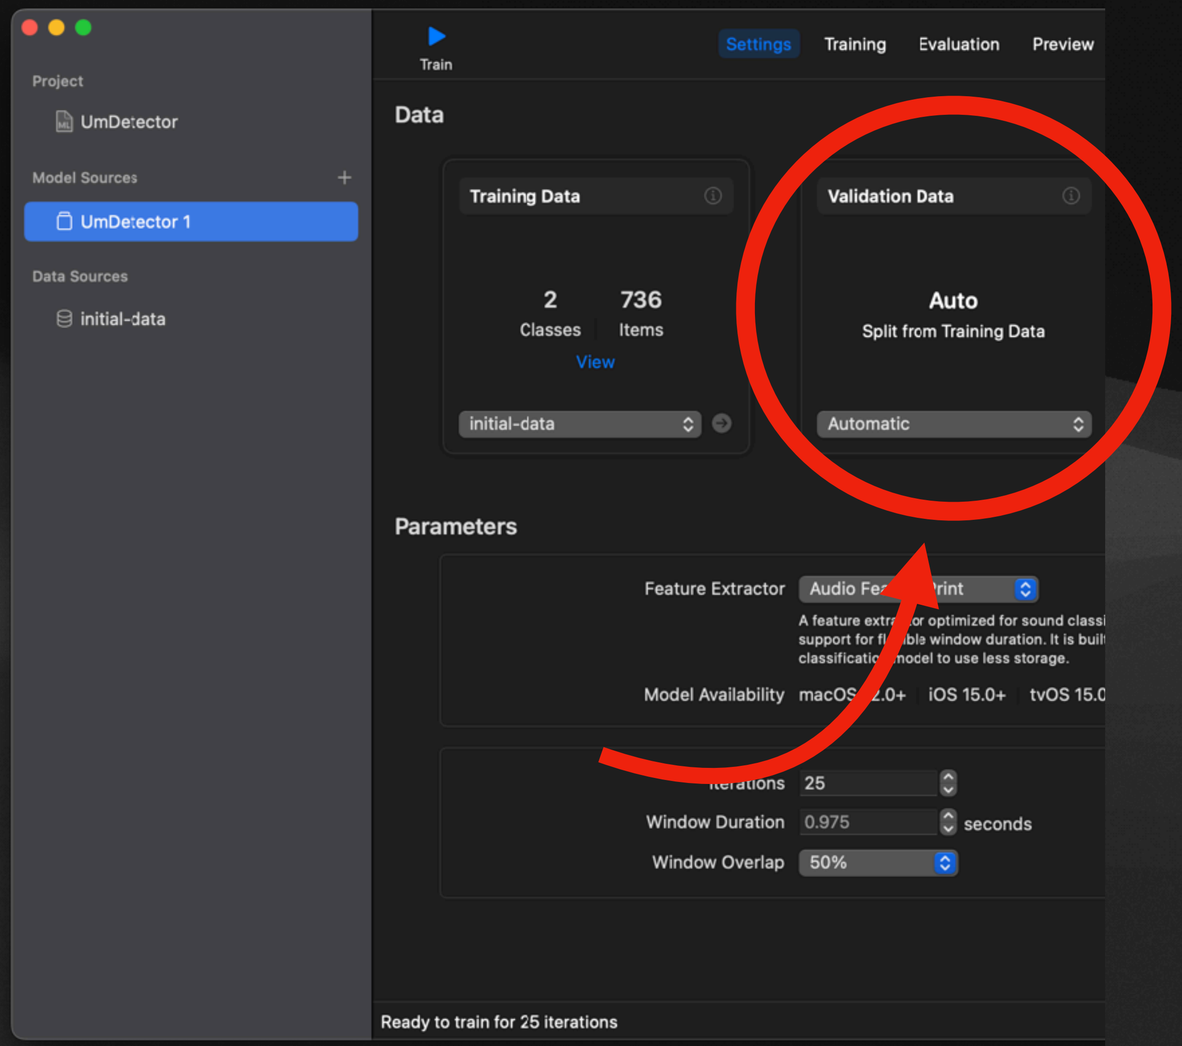

Create ML

Create a sound classification project in Create ML

Set project info

Add the dataset

Train

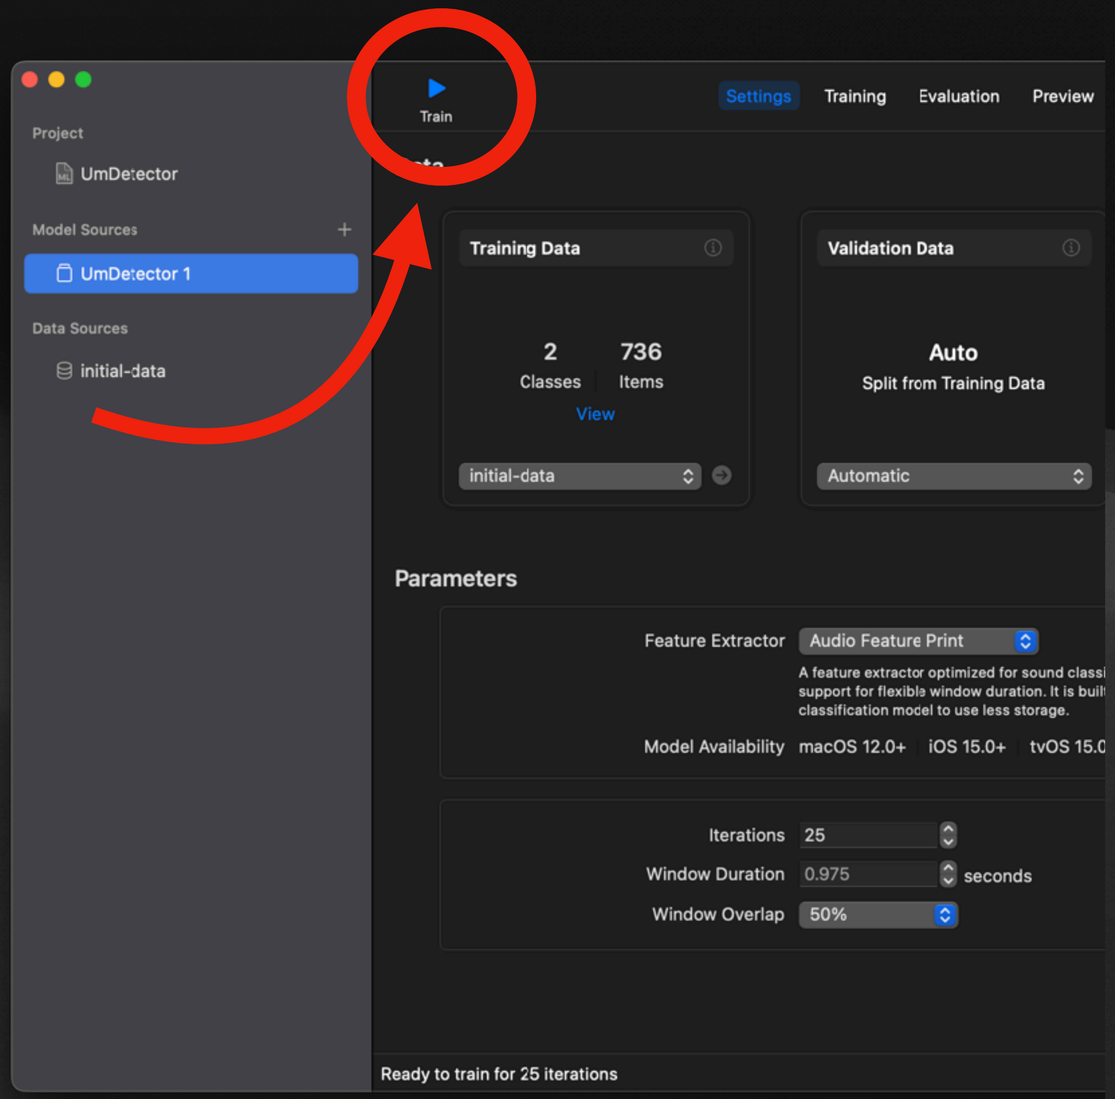

Create ML

Create a sound classification project in Create ML

Set project info

Add the dataset

Train

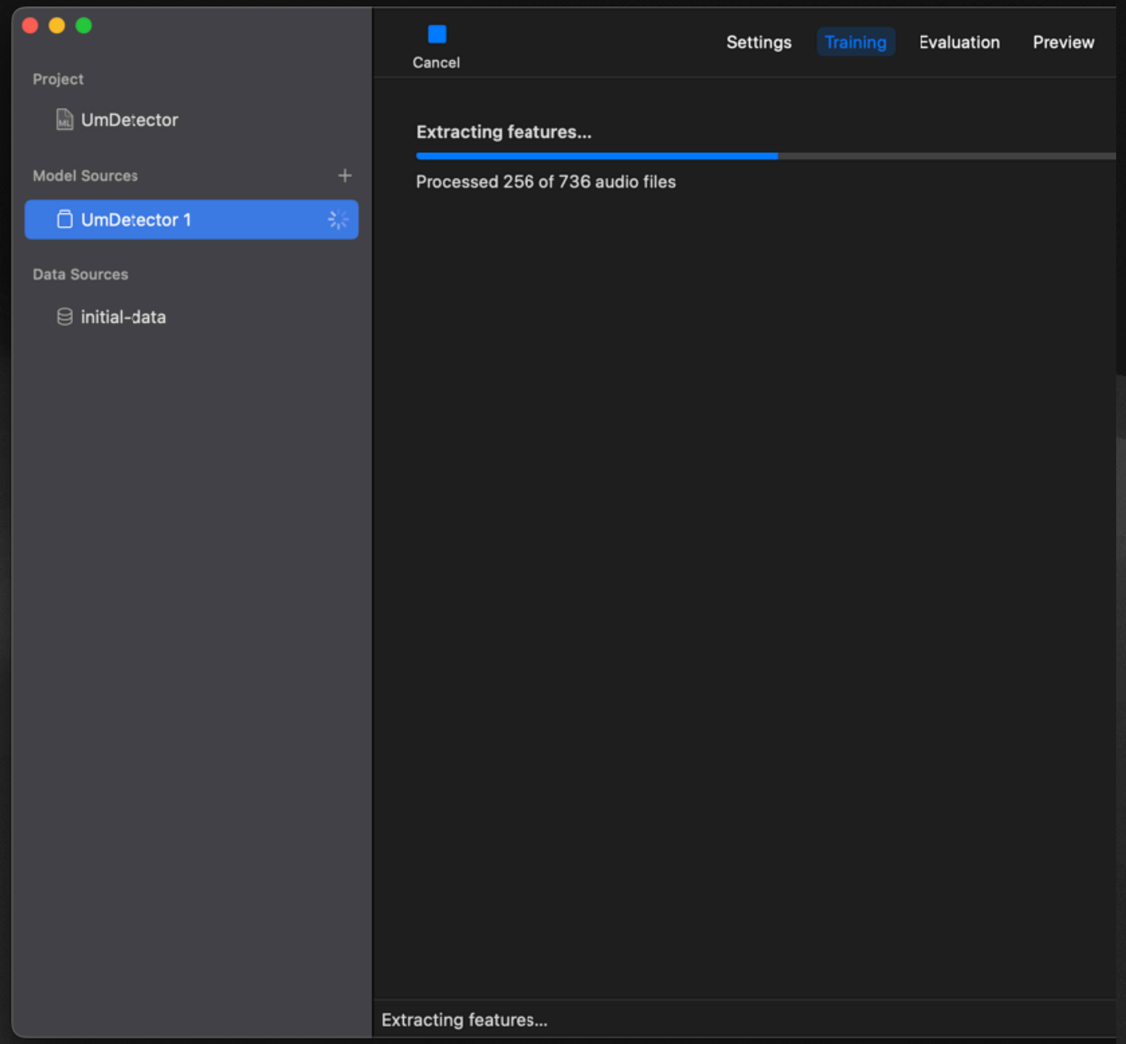

Accuracy

Training accuracy: 100%

Validation accuracy: 91.1%

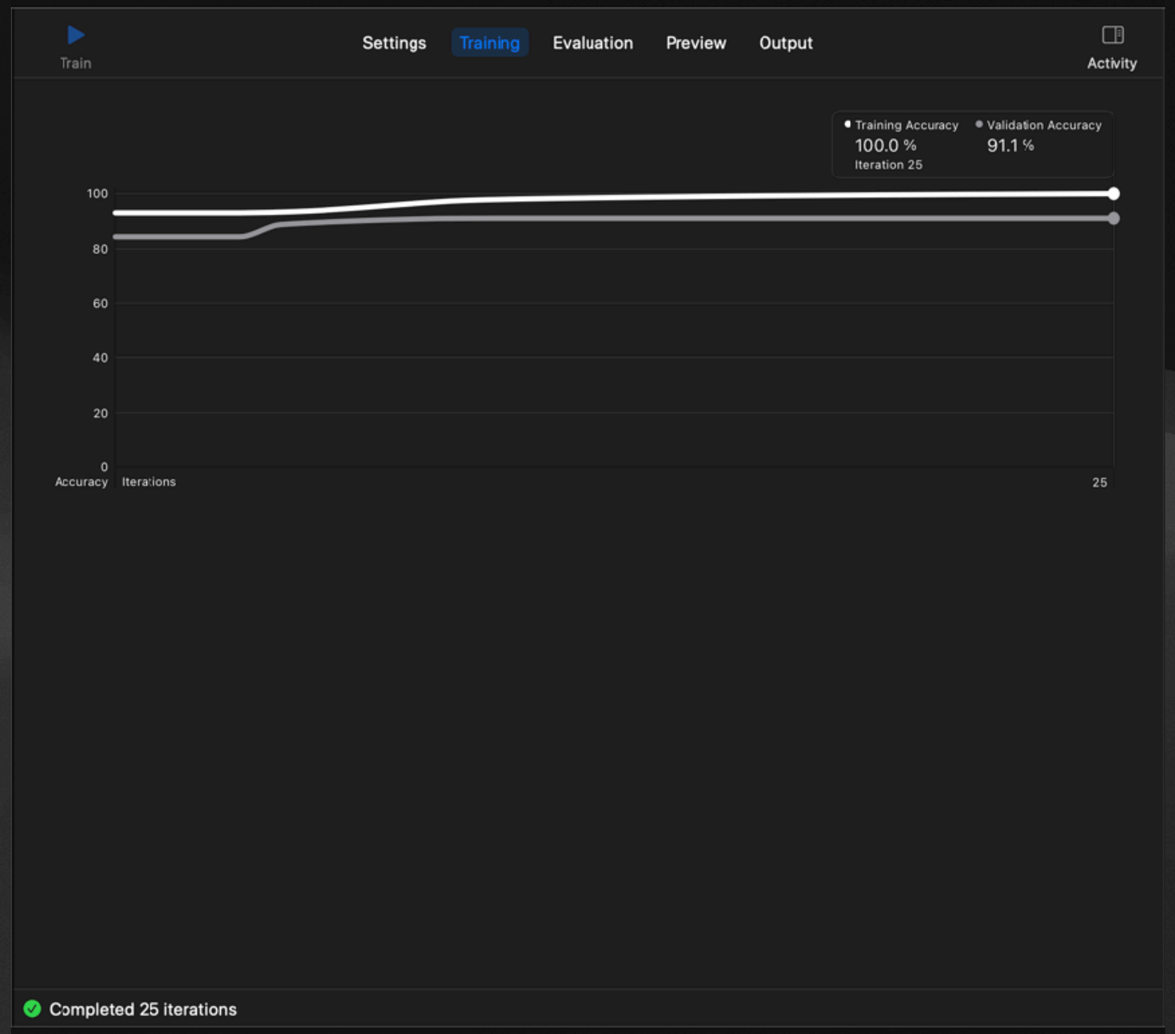

Testing - Subjective

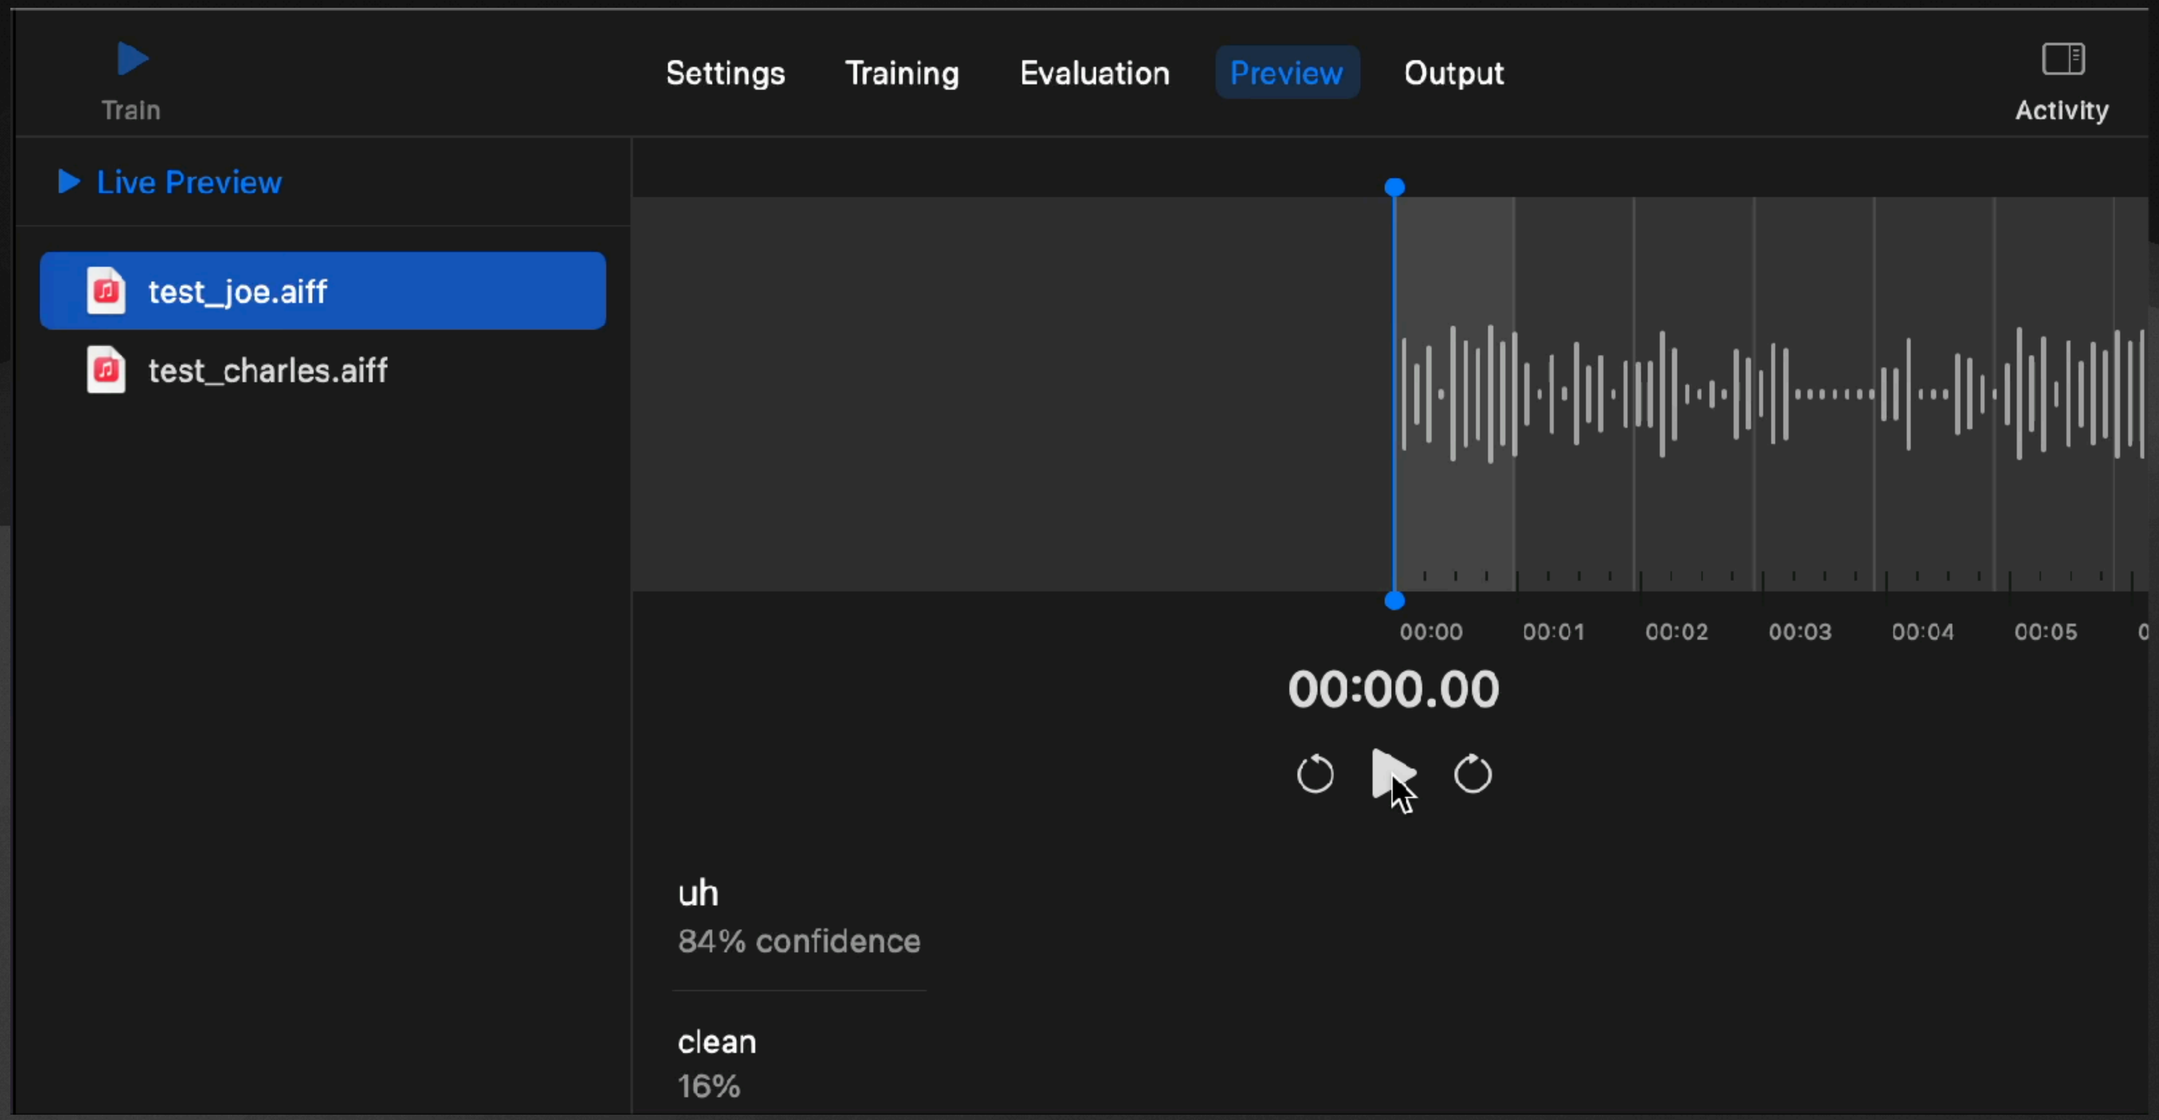

Testing - Objective

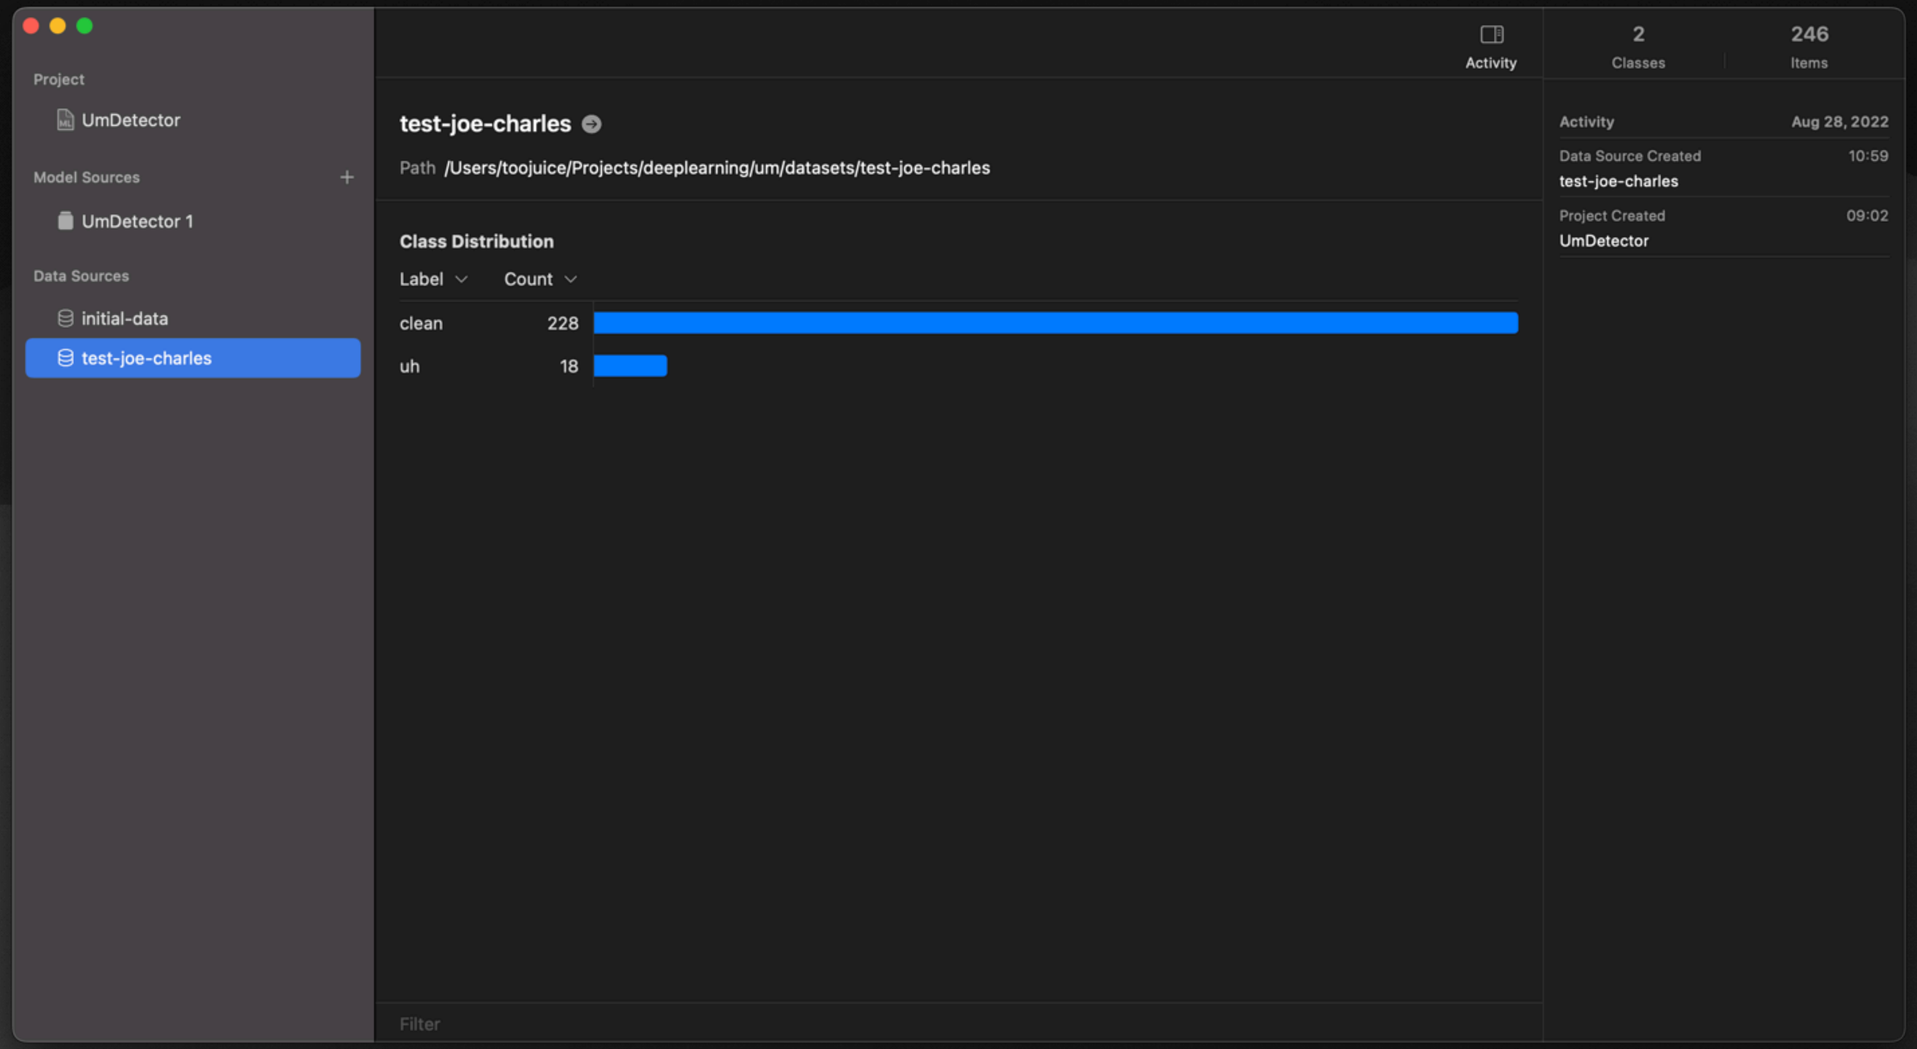

Testing - Objective

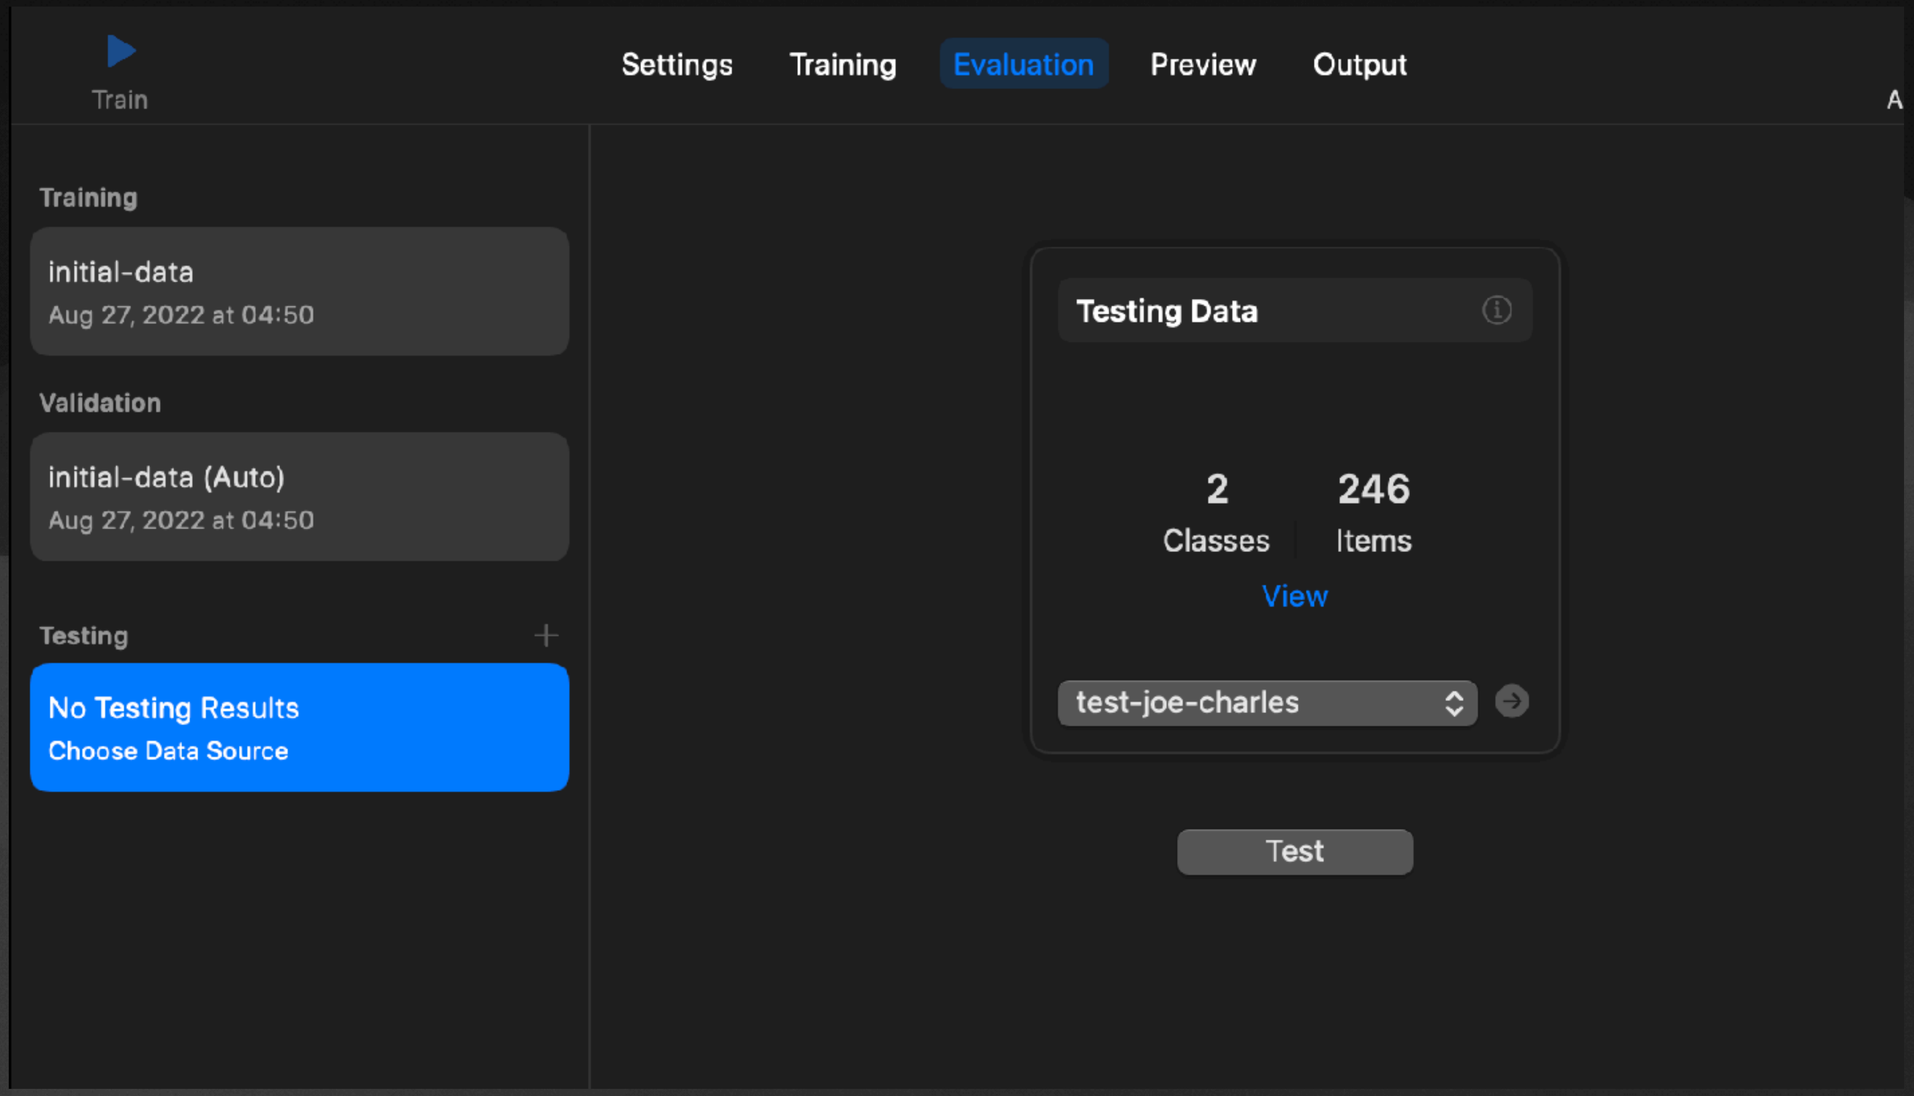

Testing - Objective

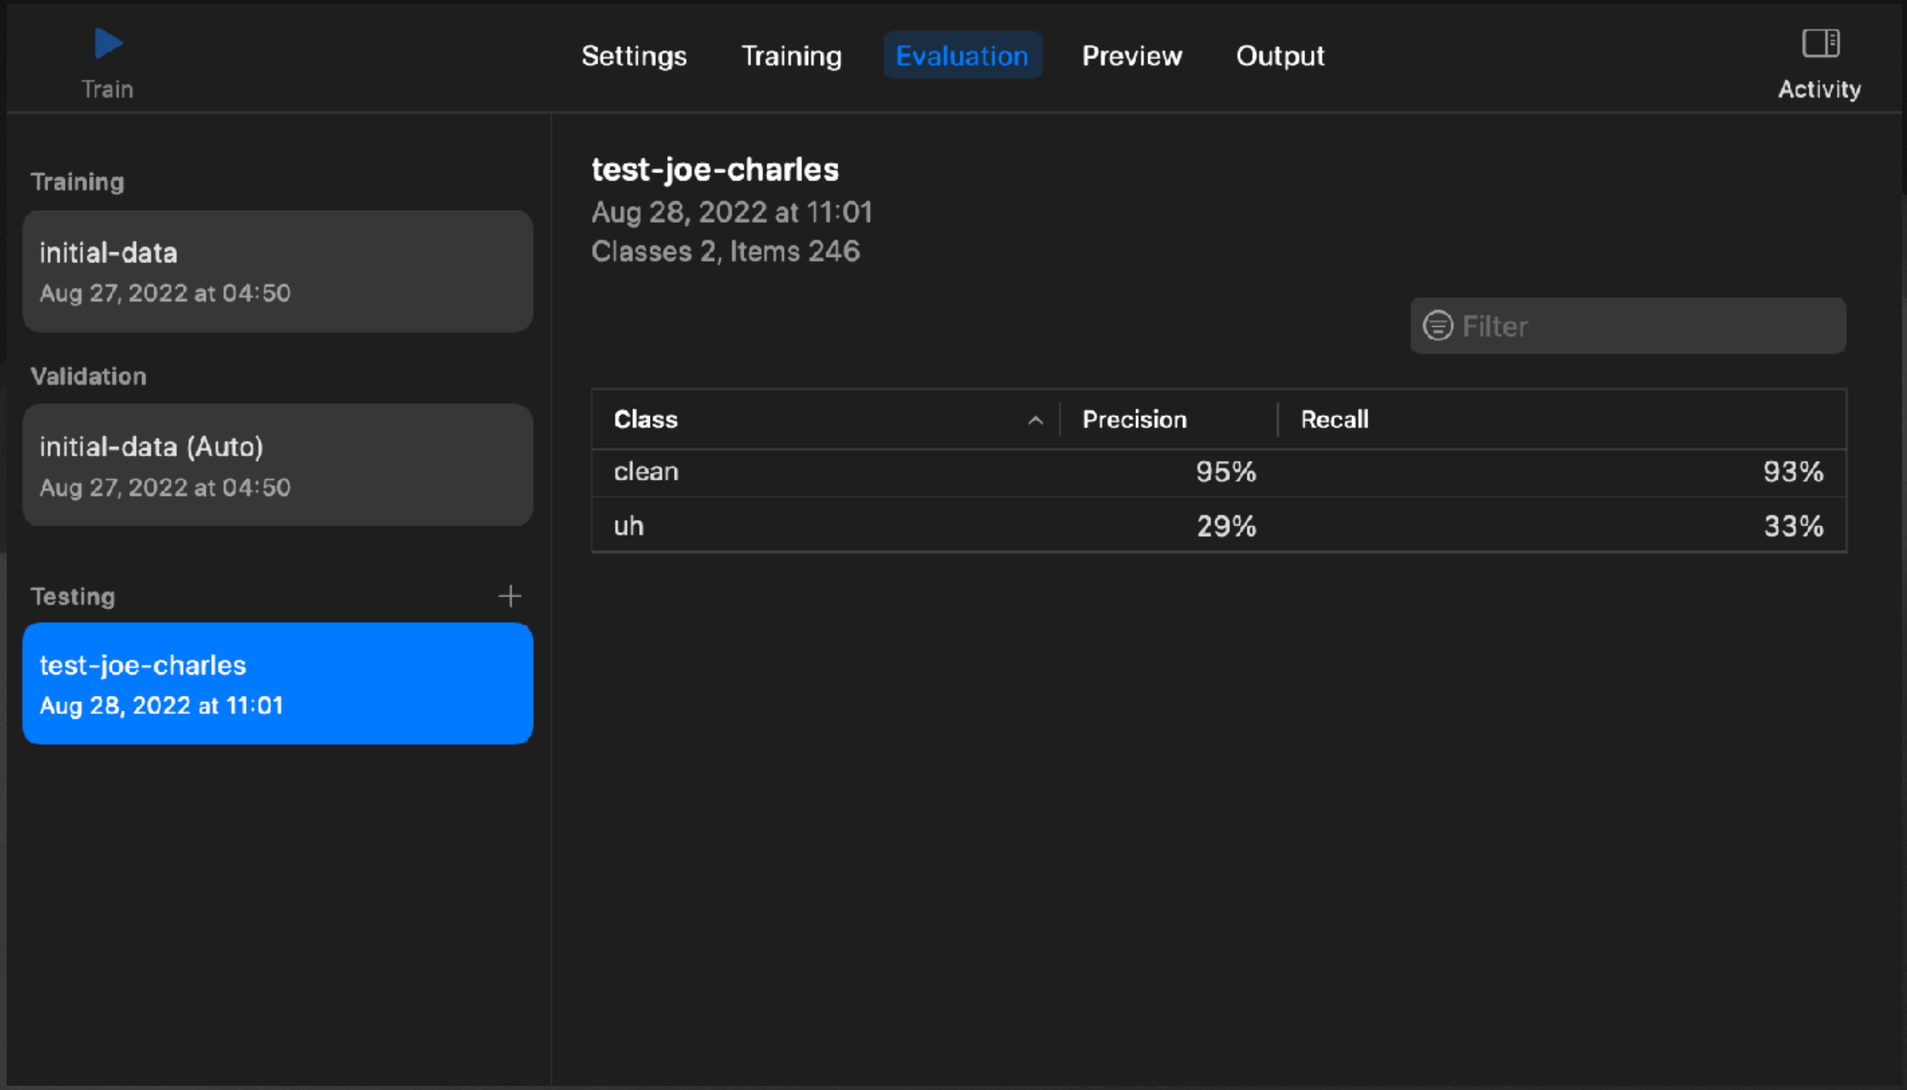

Testing - Objective

$$\frac{R_{clean} \cdot N_{clean} + R_{uh} \cdot N_{uh}}{N_{clean} + N_{uh}} = \text{Accuracy}$$

Testing - Objective

89%

IT WORKS!

sort of

Three Problems

Voice Diversity

Voice Diversity

Works well for audio recorded by Charles or Joe

Works well for me

Did not work for my wife

Did not work for my kids

Three Problems

Voice Diversity

Three Problems

Voice Diversity

Unbalanced Data

Unbalanced Data

92.5%

	Initial
Clean	681
Um	51
Total	736

Three Problems

Voice Diversity

Unbalanced Data

Labeling data is manual and time-intensive

Attempt #3 - Still Machine Learning

Augmented Data

Synthetically change our data to make the model more robust

Randomly change pitch of some samples

Randomly add noise to some samples

Ideally on the fly

Pre-baking is ok in a pinch

How do you do this?

Apple’s audio frameworks

Command line tool - sox

SOX

SOund eXchange

Many sound manipulation features

Install via Homebrew

```
> sox example.aif tmp.aif pitch 400.0 loudness -10
```

```
sox          – the command
example.aif  – input file
tmp.aif      – output file
pitch 400.0  – increase pitch by 400%
loudness -10 – decrease loudness by 10dB
```

```
> sox tmp.aif -p synth whitenoise vol 0.05 | sox -m tmp.aif - example_augmented.aif
```

sox	– the command
tmp.aif	– input file
-p synth whitenoise vol 0.05	– synthesize white noise with a volume of 0.05(?)
| sox	– pipe to the 2nd sox command
-m tmp.aif -	– merge the tmp.aif file with the output from the pipe
example_augmented.aif	– output file

https://stackoverflow.com/questions/26971240/how-do-i-run-a-terminal-command-in-a-swift-script-e-g-xcodebuild

```swift
import Foundation

func safeShell(_ command: String) throws -> String {
    let task = Process()
    let pipe = Pipe()

    task.standardOutput = pipe
    task.standardError = pipe
    task.arguments = ["-c", command]
    task.executableURL = URL(fileURLWithPath: "/bin/zsh")
    task.standardInput = nil

    try task.run()

    let data = pipe.fileHandleForReading.readDataToEndOfFile()
    let output = String(data: data, encoding: .utf8)!

    return output
}
```

https://stackoverflow.com/questions/26971240/how-do-i-run-a-terminal-command-in-a-swift-script-e-g-xcodebuild

```
do {
    print(try safeShell(“sox --help“))
} catch {
    print("\(error)")
}
```

Pitch change example 1

Pitch change example 2

Augmented Data

	Initial	Augmented
Clean	681	2046
Um	51	664
Total	736	2710

Create ML

Add new augmented dataset

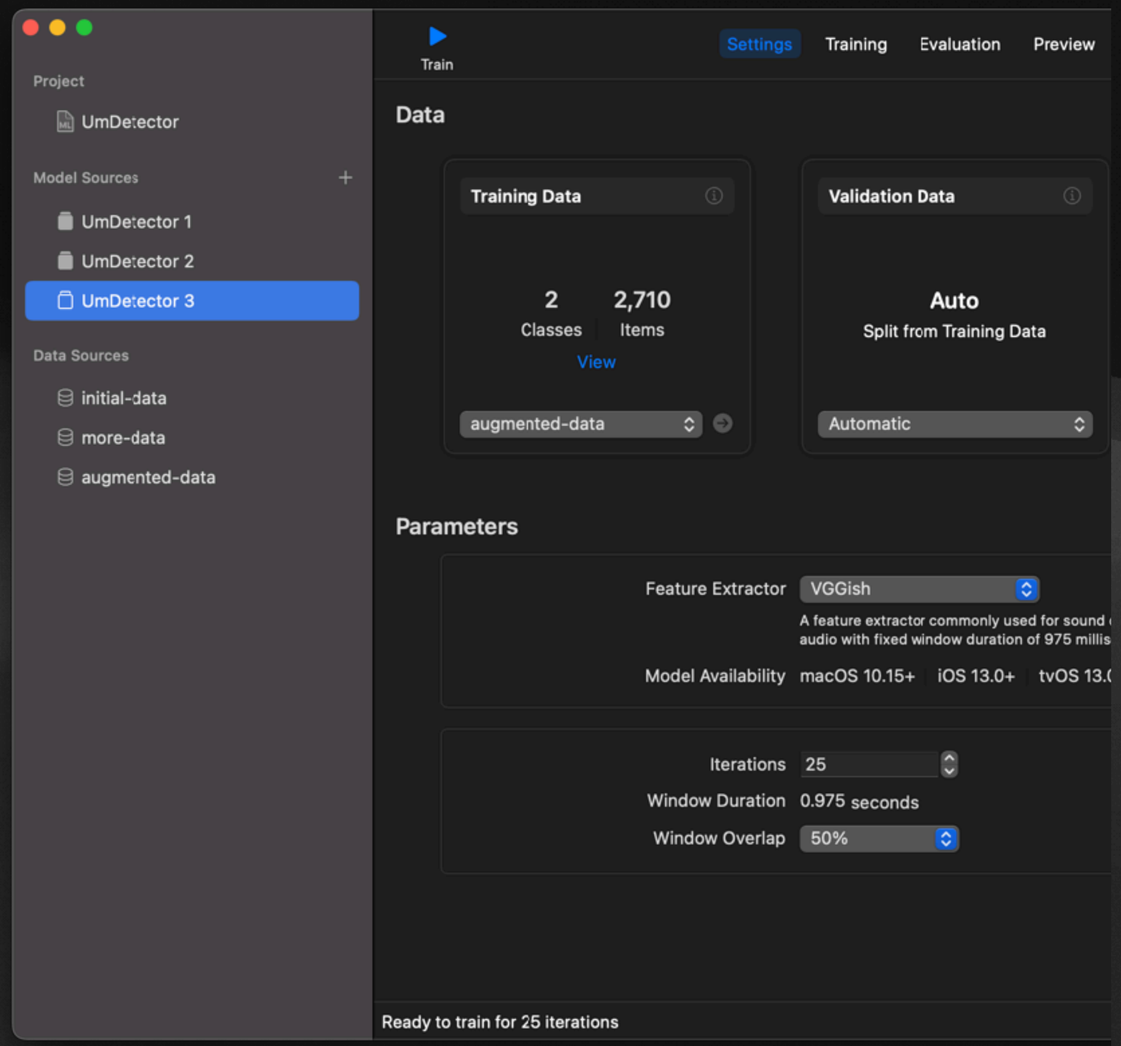

Model Change

Feature extractor

- Apple’s Audio Feature Print - 5KB
- VGGish - 4.6MB

Seemed to work better for this problem after a certain amount of data

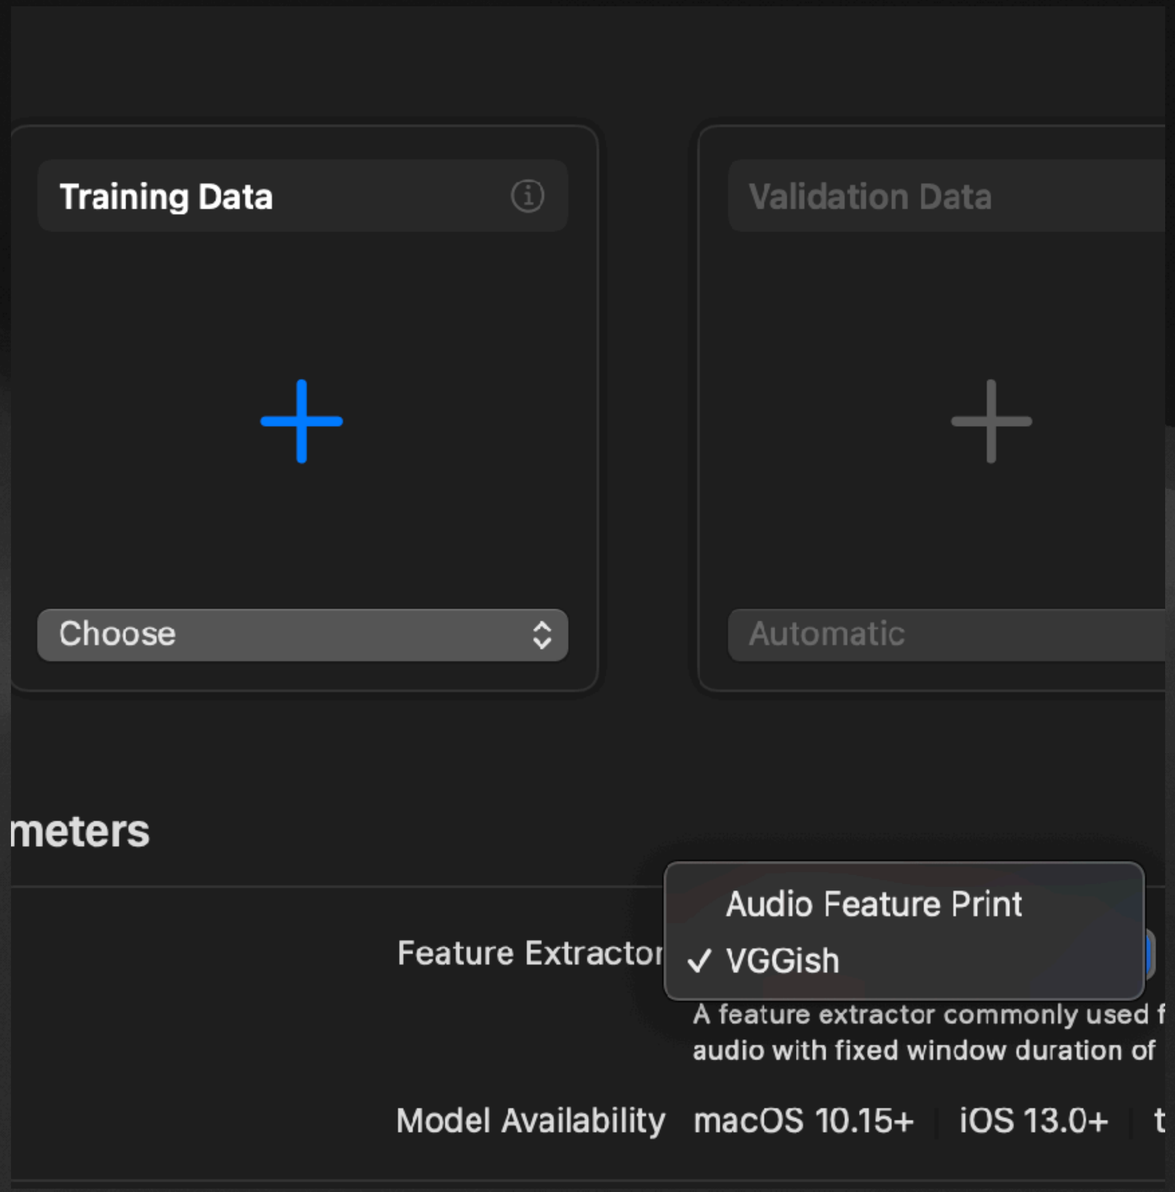

Accuracy

Training accuracy: 100.0%

Validation accuracy: 95.7%

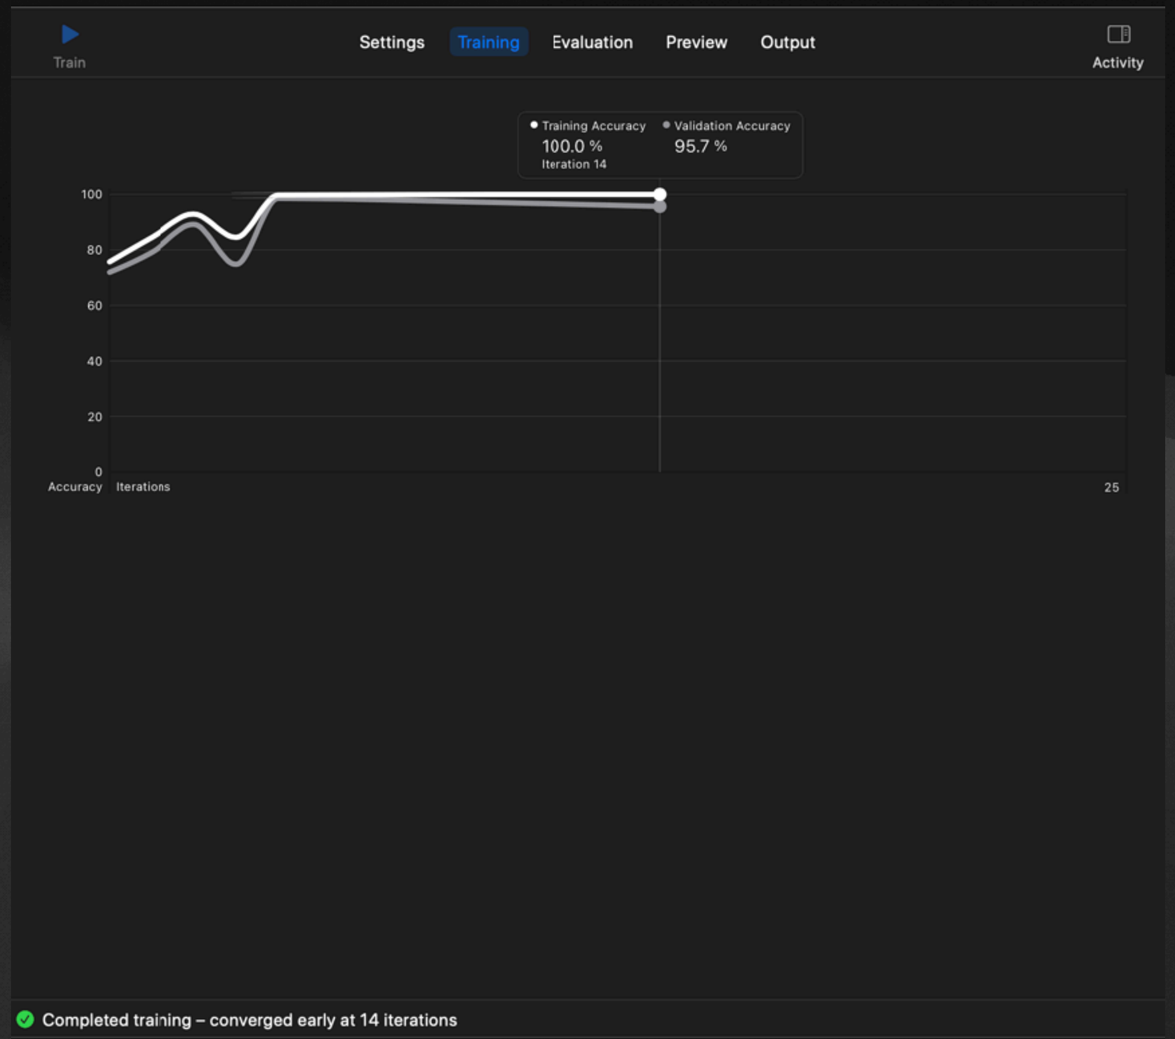

Create ML

Add new augmented dataset

Retrain

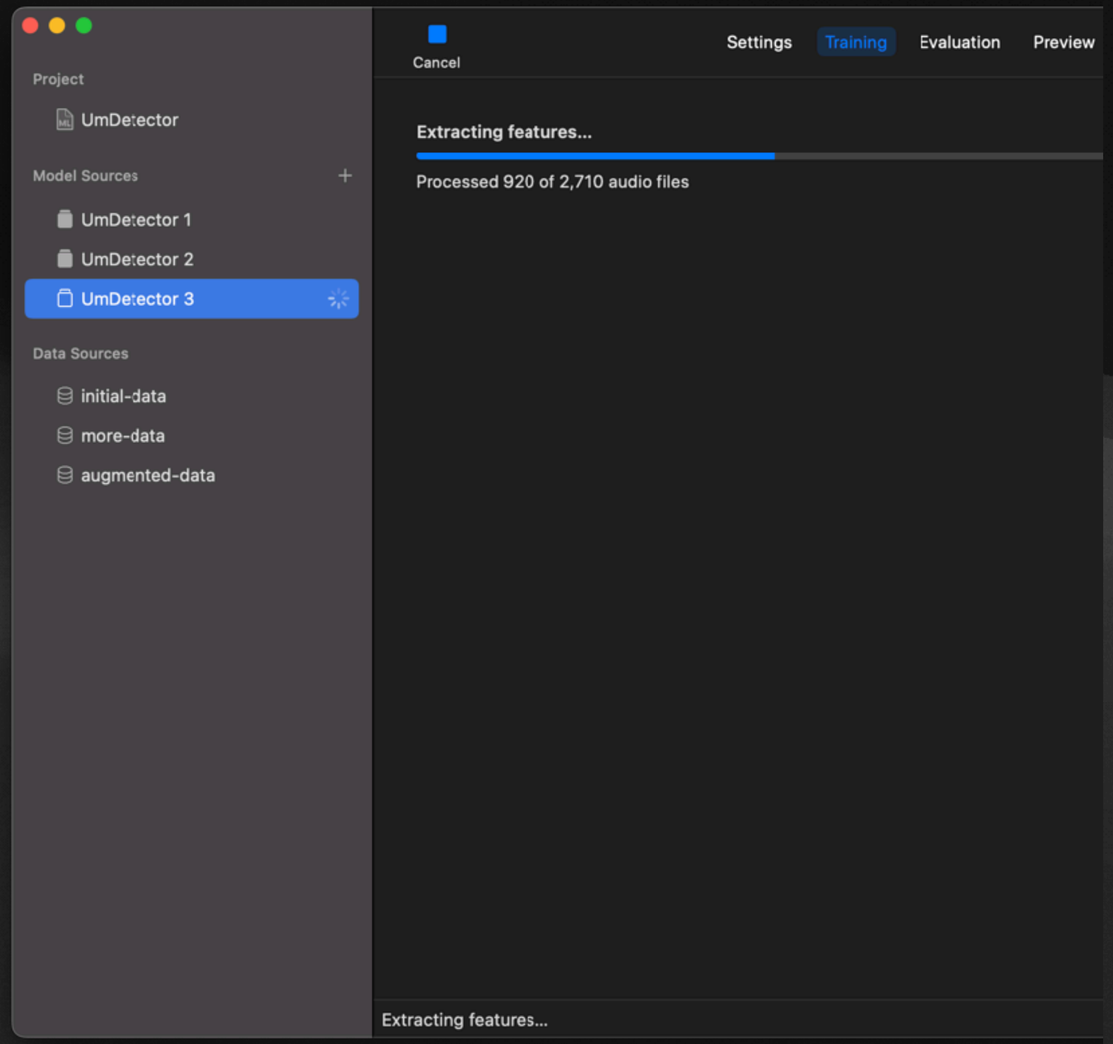

Model Change

Feature extractor

- Apple’s Audio Feature Print - 5KB
- VGGish - 4.6MB

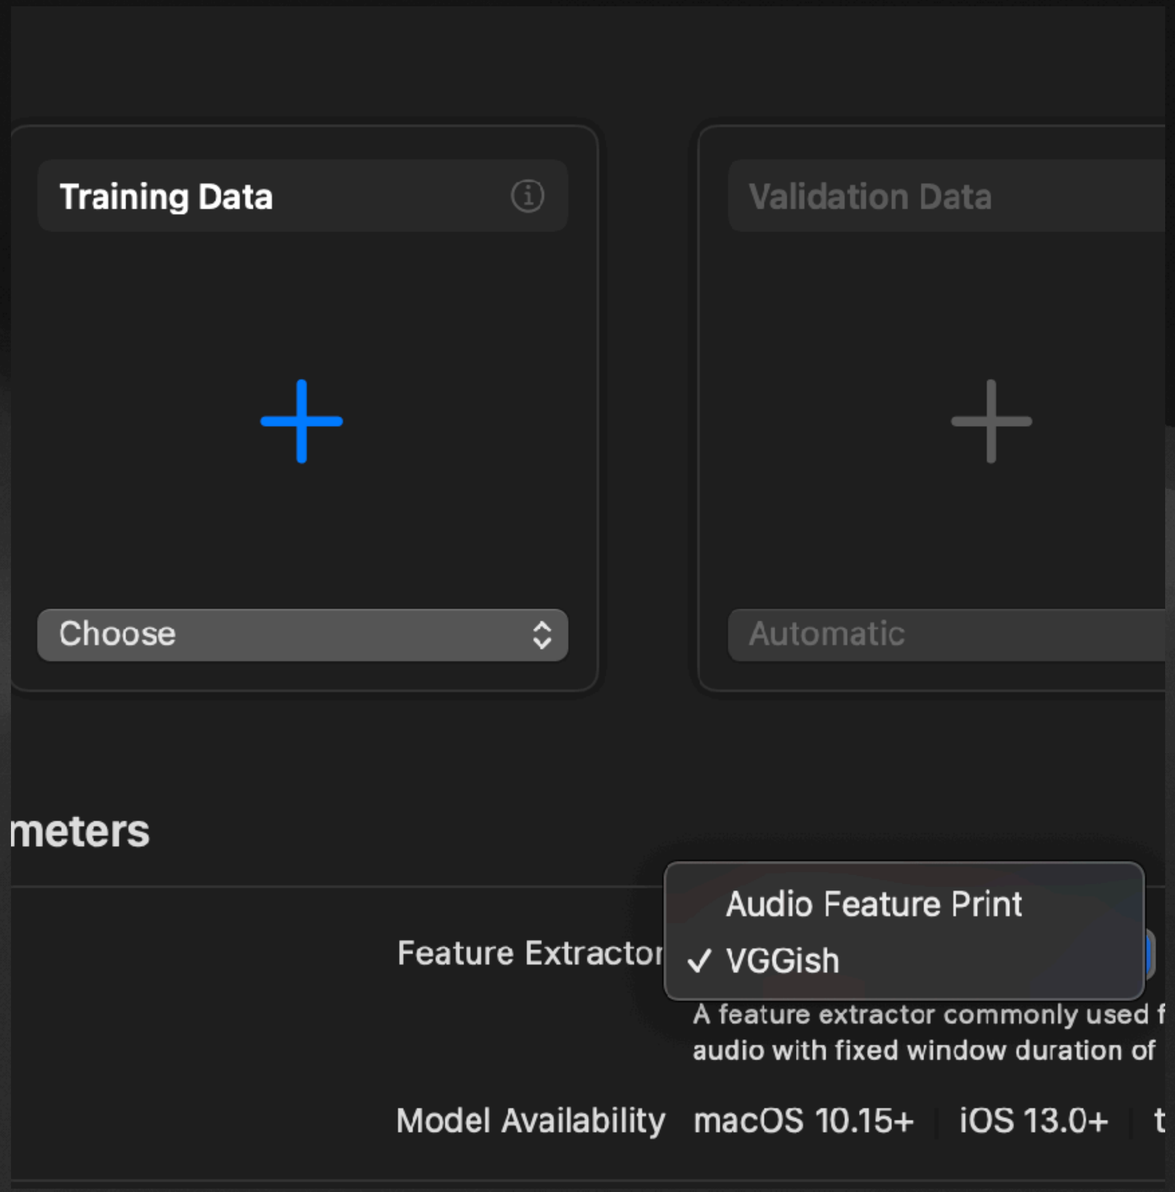

Testing - Subjective

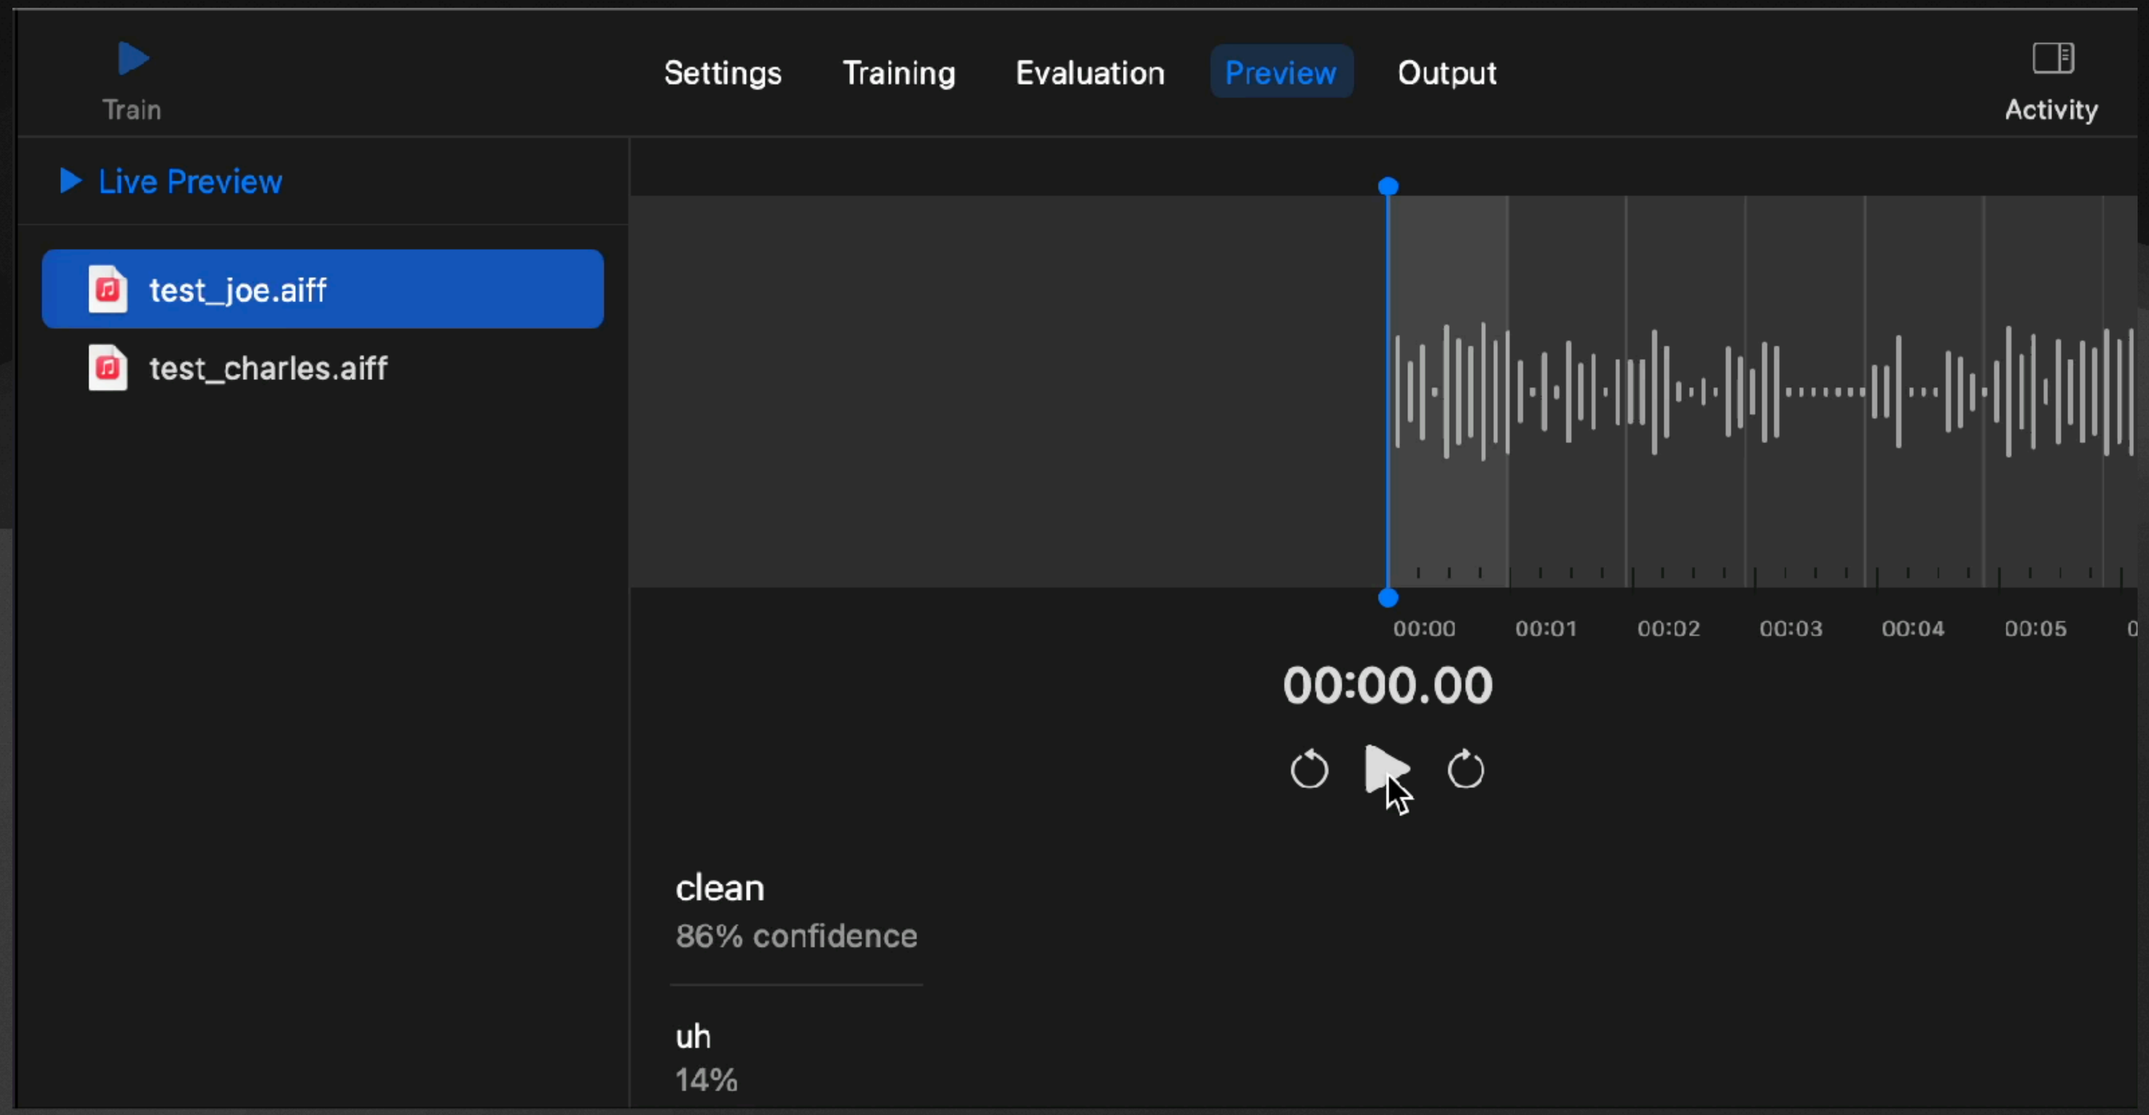

Testing - Objective

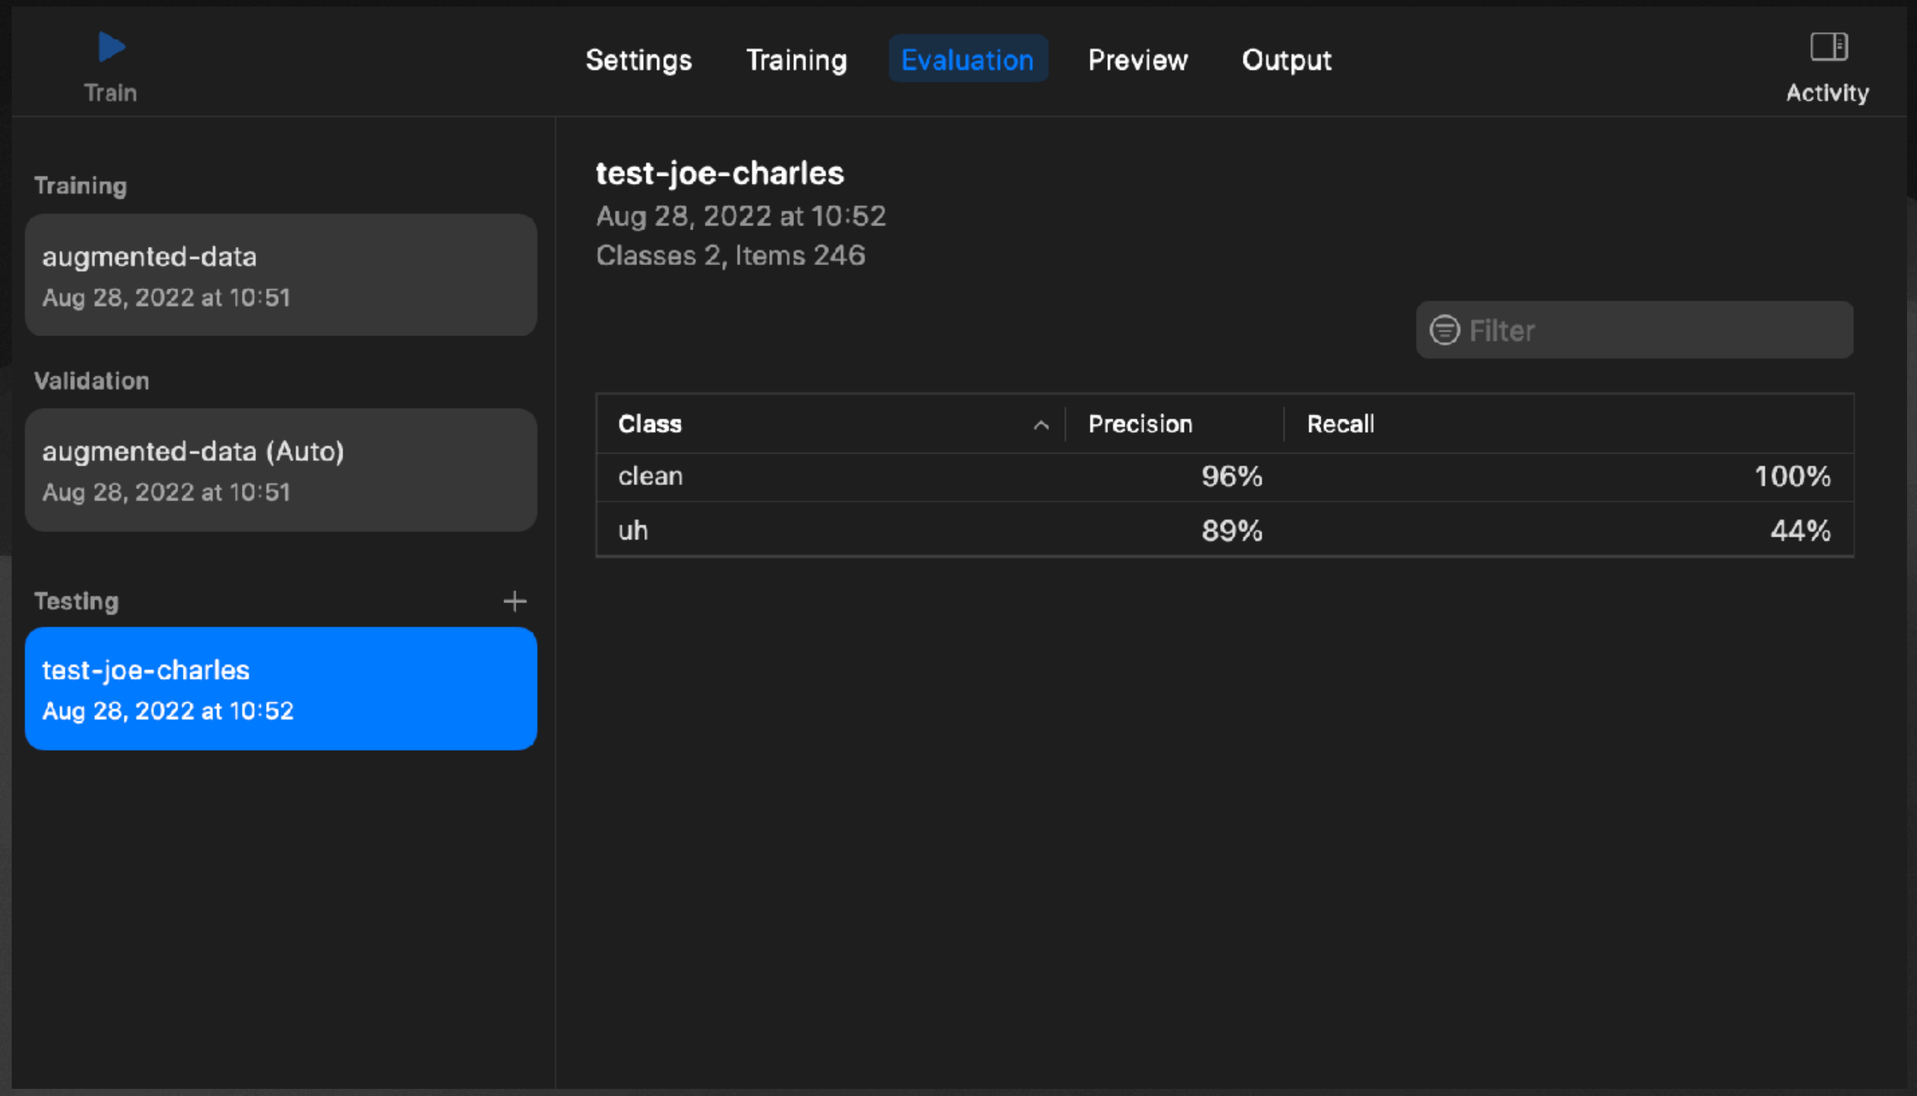

Testing - Objective

96%

What does this do for us?

Improves the accuracy and robustness of Charles and Joe recorded audio

Improves accuracy for women and children

Attempt #4 - Still More Machine Learning

MORE
DATA

LORE
DATA

More Data

Record your own audio

Find more podcasters

Open Source Datasets?

Open Source Datasets

LibriSpeech - 1000 hours of English speech

Based on public domain audiobooks

http://www.openslr.org/12

HarperValleyBank - simulated contact center calls to Harper Valley Bank

8 kHz

https://github.com/cricketclub/gridspace-stanford-harper-valley

People's Speech - licensed for academic AND commercial use

30,000 hours of transcribed audio with diverse speakers

https://mlcommons.org/en/peoples-speech/

People’s Speech

Heard about it on a podcast

Licensed for academic AND commercial use

30,000 hours of transcribed audio with diverse speakers

2TB of data

Too much data!

Active Learning

Run current model on a subset of the data

Determine which “uh”s, “um”s, and clean audio it is not super confident about

Use that to train a new model

Optionally rinse and repeat

Active Learning

Swift CLI to run inference on audio

Sort files into confidence level

Start with the least confidence files

Generated 2.5 hours of data

Only added 13 minutes to dataset

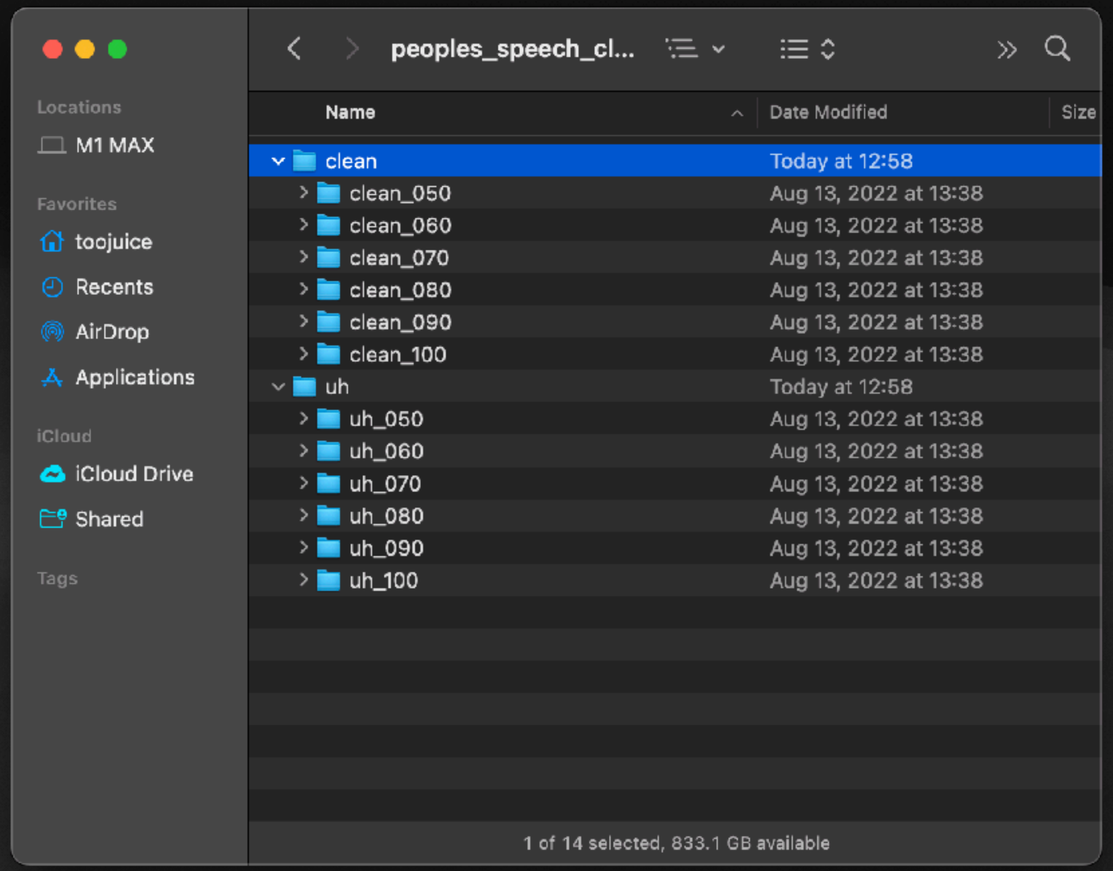

Active Learning

	Initial	Augmented	Active Learning
Clean	681	2046	2620
Um	51	664	907
Total	736	2710	3527

Create ML

Add new active learning dataset

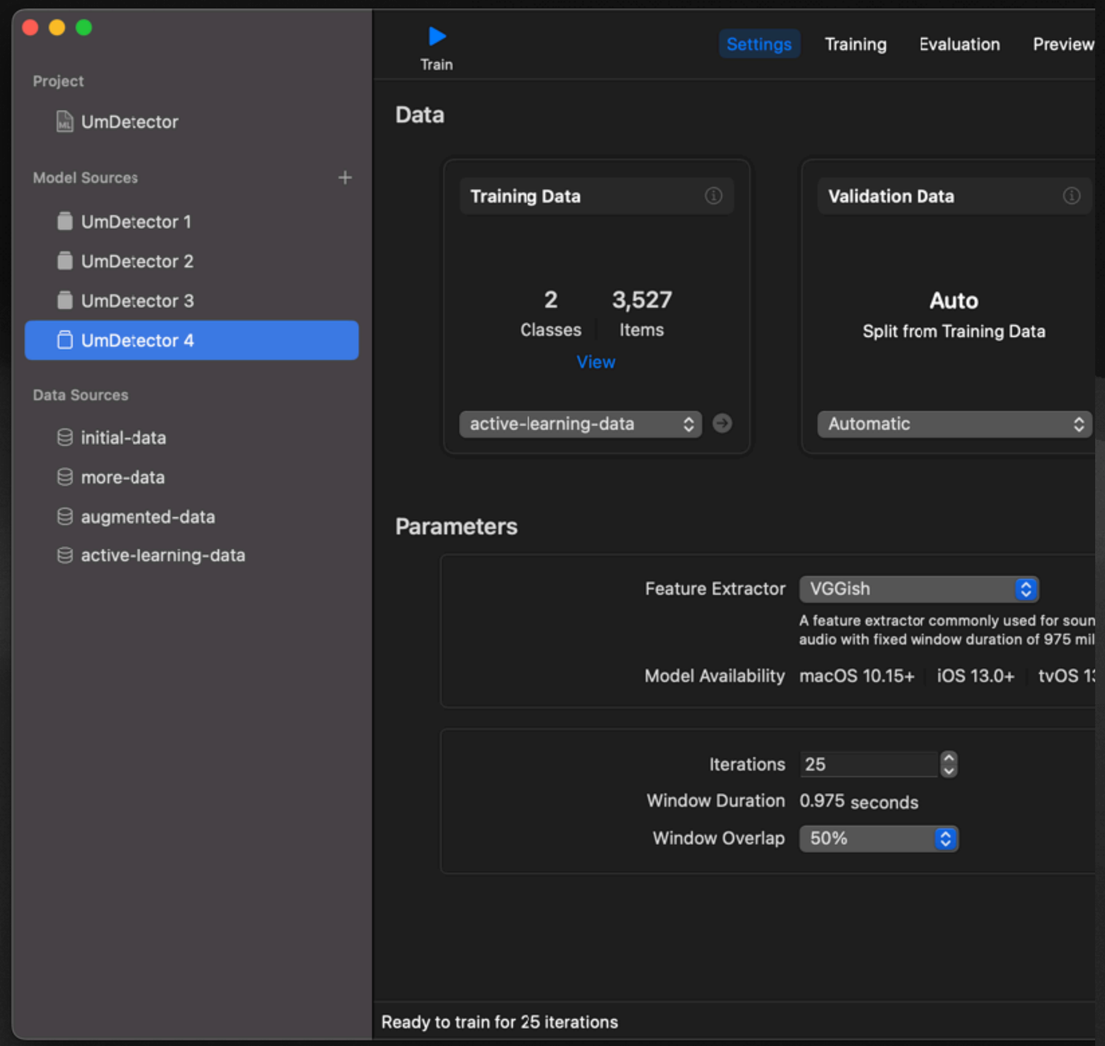

Create ML

Add new active learning dataset

Retrain

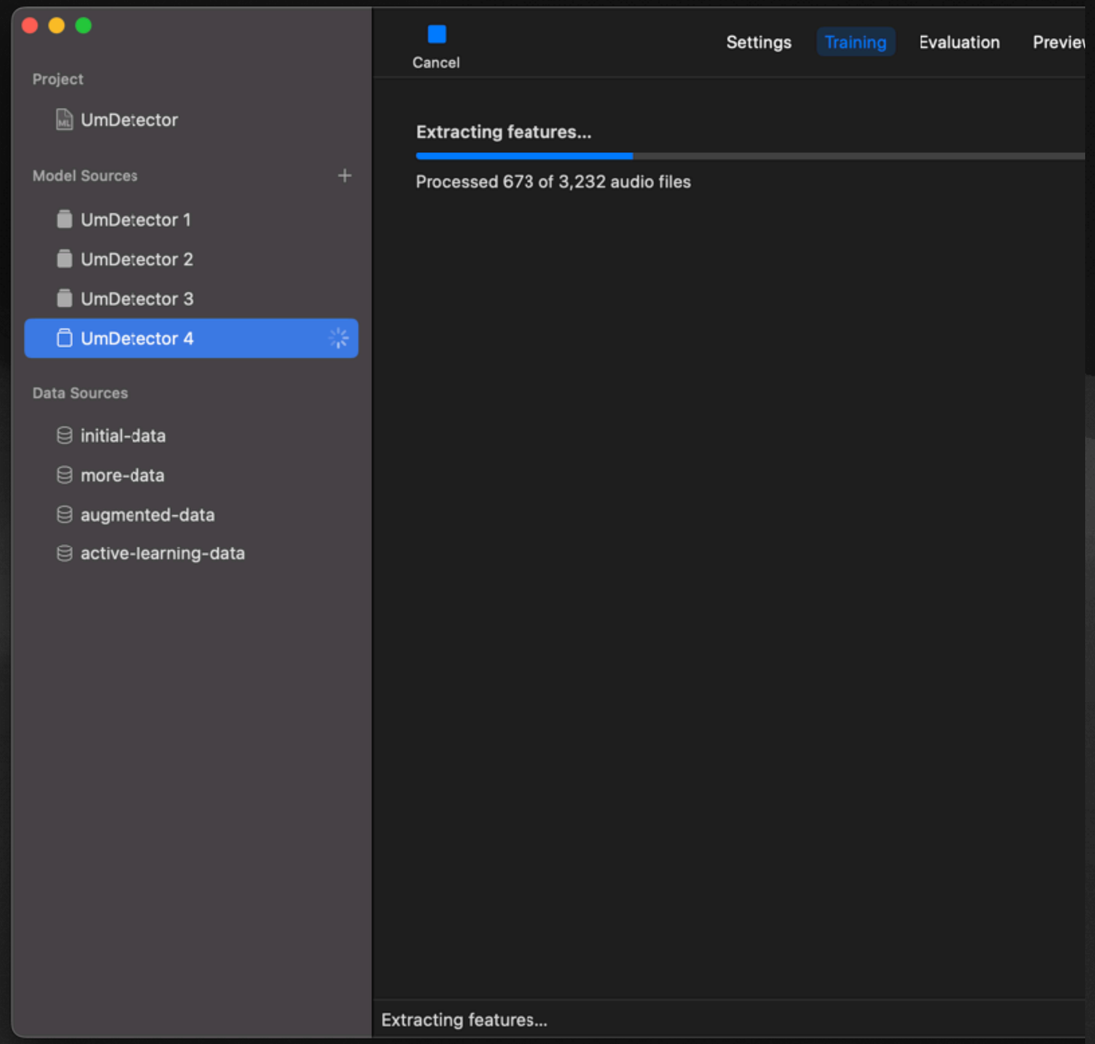

Accuracy

Training accuracy: 100.0%

Validation accuracy: 92.2%

Testing - Objective

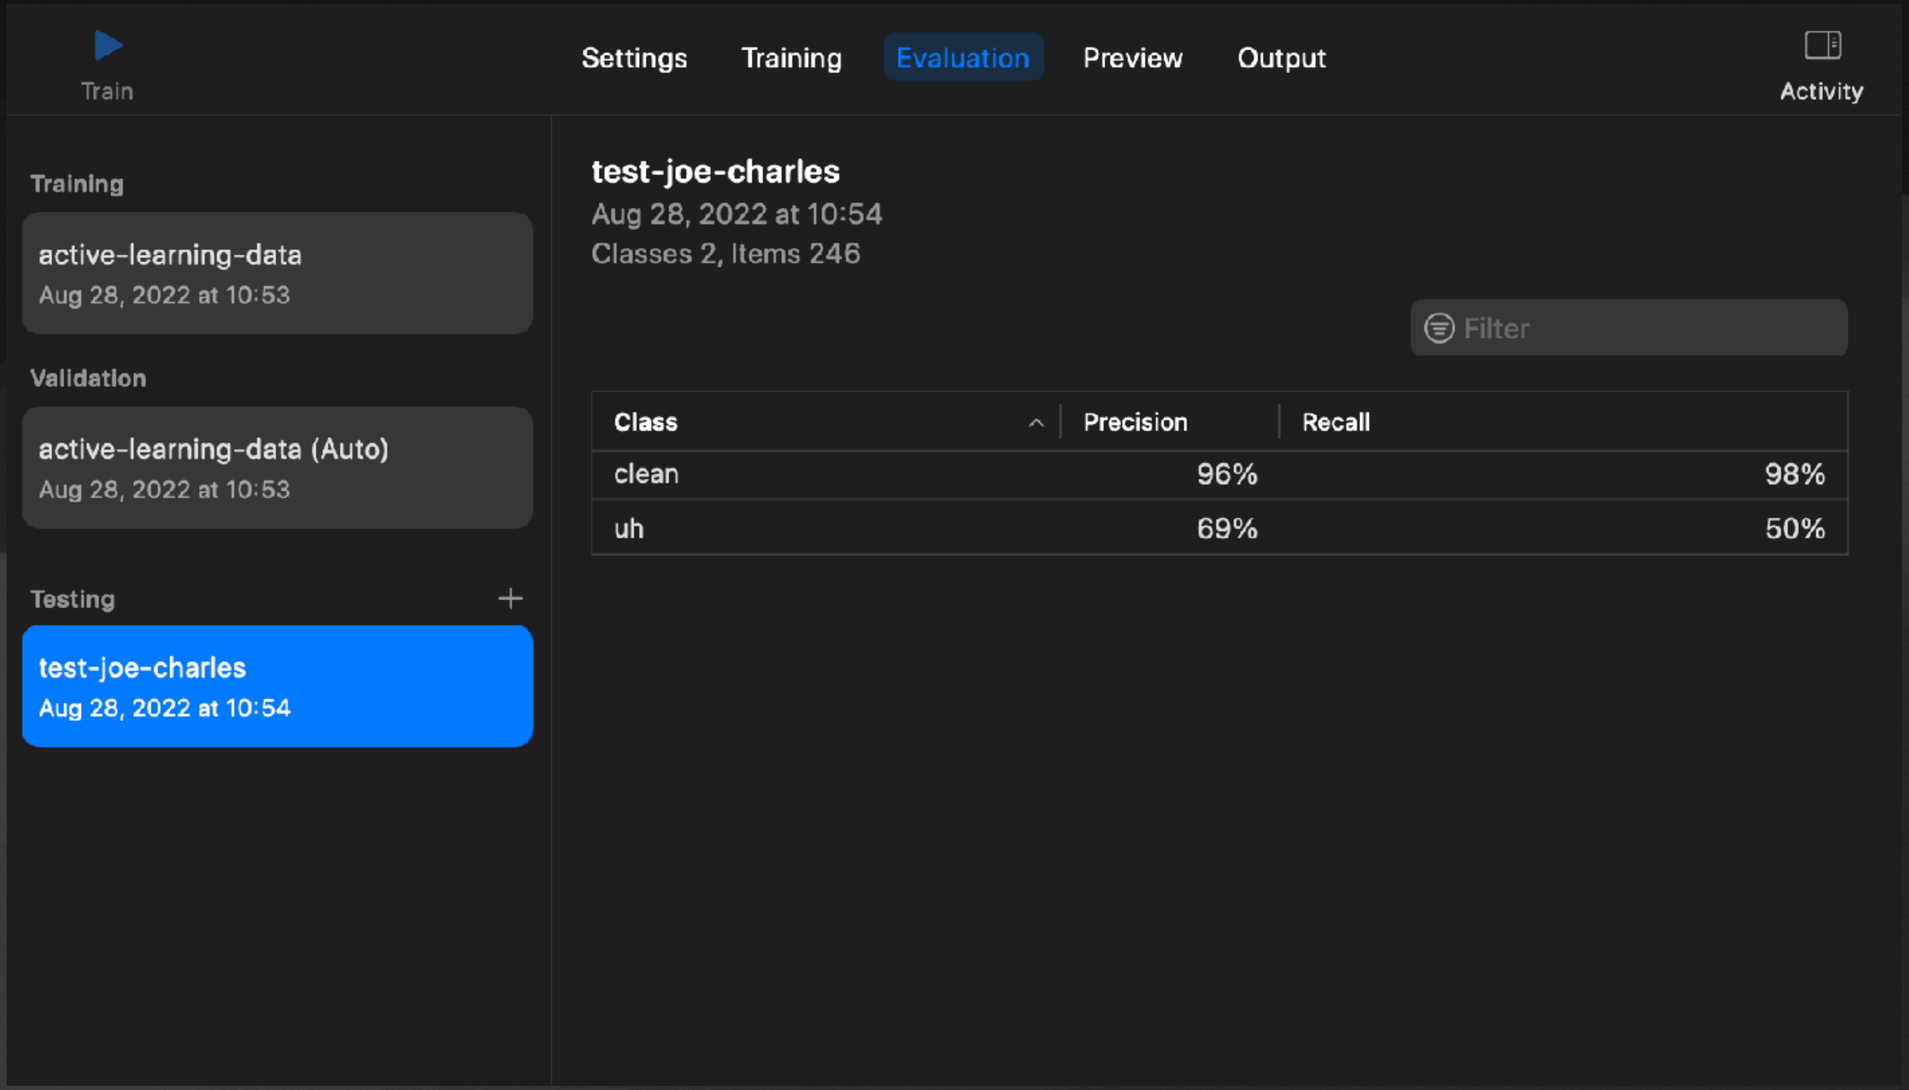

Testing - Objective

94%

Wait a minute…

Why Shouldn’t This Work?

Are we really using the right tool for the job?

Ideally, you’d have a speech recognizer that doesn’t filter out “um”s and “uh”s

This method also introduces difficulties of it’s own

Difficulties

“The” pronounced as “thuh”

“A” pronounced as “uh”

Words that begin with “un” (understand)

Different accents may be troublesome

But it worked well enough

Phew!

Summary

Speech Recognition

Collecting data

Training models using Create ML

Metrics for analyzing our models

Augmenting data to improve our models

Active learning

Lots of bad jokes

What’s left to do?

Uhhhh...

Get more data?

https://dagshub.com/yonomitt/Um-Detector

One more thing…

```
let microManager =
MicrophoneManager.shared
```

Merci

@yonomitt

https://dagshub.com/yonomitt/Um-Detector